

House Prices: Advanced Regression Techniques

Abstract:

House Price prediction is a very popular dataset for data science competition. In this dataset 79 explanatory variables describing (almost) every aspect of residential homes in Ames and Iowa. This competition challenges competitor to predict the final price of each home.

In this report my main focus is how artificial neural network performs for this kind of problems and how to improve performance of the prediction using artificial neural network. So my elaboration on that section will be much more detailed. I have divided my work in four part and they are

- **Data processing** where I have visualized, cleaned, handled missing data, carefully modified , removed and merged some features.
- **Testing multiple model** In this part I have used gradient boosting, decision tree, random forest regression , lasso and Artificial neural network on my pre processed data.
- **Artificial neural network implementation** In this section I have implemented ANN , performed parameter tuning, training, used grid search inside training and validate test score.
- **Cross Validation** In this part I have used k fold cross validation on my artificial neural network model to make sure if the Data is actually independent and to fine tune few parameters on whole dataset if the cross validation score is not same as validation score.
- **Ensemble learning** I have used bagging method for this section to improve my kaggle score.

Score:

Best Score : 0.12192 (using Ensemble Learning)

output.csv 18 hours ago by Navid Using ['ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None' 'ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16' 'ANN_lr0.05_beta0.005-0.1-0.05-0.0_hidden8-32-16-8' 'ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2' 'Random Forest Regressor' 'Xgboost' 'Lasso'] * [.15,,1,,1,,05,,0,,2,,4]	0.12192	<input type="checkbox"/>
--	---------	--------------------------

Best score without Ensemble : 0.12324 (ANN only)

output.csv a month ago by navid ann	0.12324	<input checked="" type="checkbox"/>
---	---------	-------------------------------------

Imports:

Gpu testing

```
In [4]: import tensorflow as tf
device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')
print('Found GPU at: {}'.format(device_name))

Found GPU at: /device:GPU:0
```

```
In [5]: import tensorflow as tf
import numpy as np
import pandas as pd
from sklearn.ensemble import RandomForestRegressor
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler
from IPython.display import Image
from sklearn.preprocessing import normalize, MinMaxScaler
import matplotlib.pyplot as plt
from sklearn.utils import shuffle
# %matplotlib widget
%matplotlib inline
```

Data Pre-processing

Load Data

```
In [6]: train = pd.read_csv('train.csv').select_dtypes(exclude=['object'])
test = pd.read_csv('test.csv').select_dtypes(exclude=['object'])

#look into datatypes of the file
print("data types count")
train.dtypes.groupby(train.dtypes).count()
```

data types count

```
Out[6]: int64      35
float64      3
object      43
dtype: int64
```

Looking into data

```
In [7]: print('show sample')
pd.set_option('display.max_column', None)
train.head()
```

show sample

Out[7]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	Condition2	BldgType	...
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	Norm	1Fam	...
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	FR2	Gtl	Veenker	Feedr	Norm	1Fam	...
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	Norm	1Fam	...
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	Corner	Gtl	Crawfor	Norm	Norm	1Fam	...
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	FR2	Gtl	NoRidge	Norm	Norm	1Fam	...

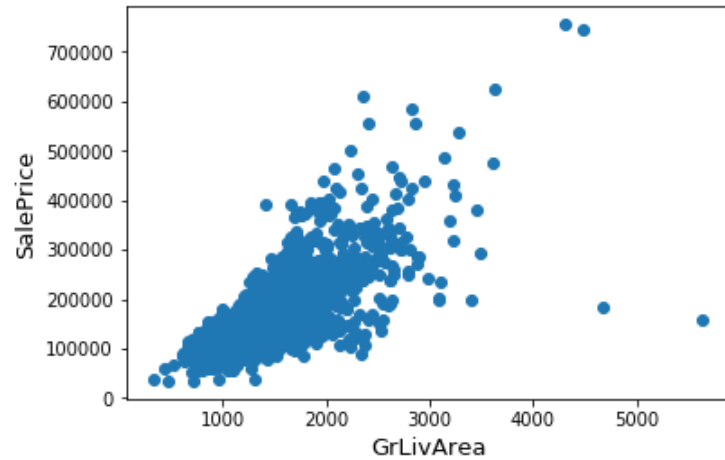
```
In [8]: print('description of data')
train.describe()
```

description of data

Out[8]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1	BsmtFinSF2	BsmtUnfSF	Tot
count	1460.000000	1460.000000	1201.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1452.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	730.500000	56.897260	70.049958	10516.828082	6.099315	5.575342	1971.267808	1984.865753	103.685262	443.639726	46.549315	567.240411	10000.000000
std	421.610009	42.300571	24.284752	9981.264932	1.382997	1.112799	30.202904	20.645407	181.066207	456.098091	161.319273	441.866955	4000.000000
min	1.000000	20.000000	21.000000	1300.000000	1.000000	1.000000	1872.000000	1950.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	365.750000	20.000000	59.000000	7553.500000	5.000000	5.000000	1954.000000	1967.000000	0.000000	0.000000	0.000000	223.000000	7000.000000
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	5.000000	1973.000000	1994.000000	0.000000	383.500000	0.000000	477.500000	9000.000000
75%	1095.250000	70.000000	80.000000	11601.500000	7.000000	6.000000	2000.000000	2004.000000	166.000000	712.250000	0.000000	808.000000	12000.000000
max	1460.000000	190.000000	313.000000	215245.000000	10.000000	9.000000	2010.000000	2010.000000	1600.000000	5644.000000	1474.000000	2336.000000	61000.000000

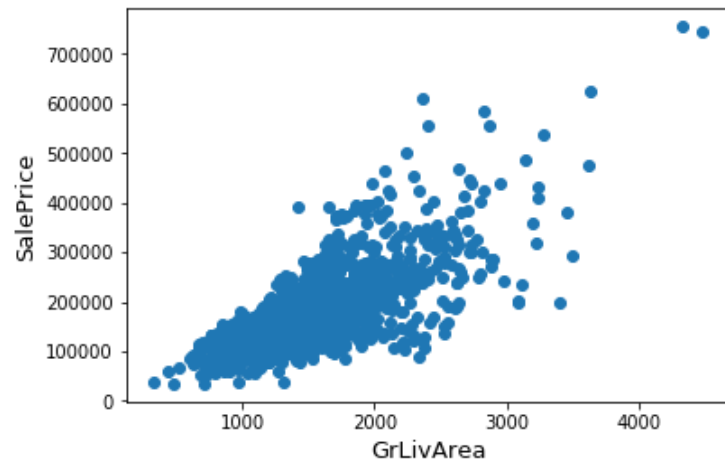
```
In [9]: fig, ax = plt.subplots()
ax.scatter(x = train['GrLivArea'], y = train['SalePrice'])
plt.ylabel('SalePrice', fontsize=13)
plt.xlabel('GrLivArea', fontsize=13)
plt.show()
```



There are a few houses with more than 4000 sq ft living area that are outliers, so we drop them from the training data.

```
In [10]: train.drop(train[ (train["GrLivArea"] > 4000) & (train['SalePrice']<400000) ].index, inplace=True)
```

```
In [11]: #Check the graph again
fig, ax = plt.subplots()
ax.scatter(train['GrLivArea'], train['SalePrice'])
plt.ylabel('SalePrice', fontsize=13)
plt.xlabel('GrLivArea', fontsize=13)
plt.show()
```



RMSE

```
In [12]: def rmse(y_true, y_pred):
          return np.sqrt(mean_squared_error(y_true, y_pred))
```

Imputing missing data

```
In [13]: lot_frontage_by_neighborhood = train["LotFrontage"].groupby(train["Neighborhood"])
```

```
In [14]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()

def factorize(df, factor_df, column, fill_na=None):
    factor_df[column] = df[column]
    if fill_na is not None:
        factor_df[column].fillna(fill_na, inplace=True)
    le.fit(factor_df[column].unique())
    factor_df[column] = le.transform(factor_df[column])
    return factor_df
```

common data processing:

In this part we have label encoded some of the columns because some features are ordinal. I have replaced some null value with zero because in those case they probably meant that it may not exist . Finally I have merged some of the features to get a better feature.

Befor starting following block its important to understand which feature means what so that describing my work would be easier

- SalePrice - the property's sale price in dollars. This is the target variable that you're trying to predict.
- MSSubClass: The building class
- MSZoning: The general zoning classification
- LotFrontage: Linear feet of street connected to property
- LotArea: Lot size in square feet
- Street: Type of road access
- Alley: Type of alley access
- LotShape: General shape of property
- LandContour: Flatness of the property
- Utilities: Type of utilities available
- LotConfig: Lot configuration
- LandSlope: Slope of property
- Neighborhood: Physical locations within Ames city limits
- Condition1: Proximity to main road or railroad
- Condition2: Proximity to main road or railroad (if a second is present)
- BldgType: Type of dwelling
- HouseStyle: Style of dwelling
- OverallQual: Overall material and finish quality
- OverallCond: Overall condition rating
- YearBuilt: Original construction date
- YearRemodAdd: Remodel date
- RoofStyle: Type of roof
- RoofMatl: Roof material
- Exterior1st: Exterior covering on house
- Exterior2nd: Exterior covering on house (if more than one material)
- MasVnrType: Masonry veneer type
- MasVnrArea: Masonry veneer area in square feet
- ExterQual: Exterior material quality
- ExterCond: Present condition of the material on the exterior
- Foundation: Type of foundation
- BsmtQual: Height of the basement
- BsmtCond: General condition of the basement
- BsmtExposure: Walkout or garden level basement walls
- BsmtFinType1: Quality of basement finished area
- BsmtFinSF1: Type 1 finished square feet
- BsmtFinType2: Quality of second finished area (if present)
- BsmtFinSF2: Type 2 finished square feet
- BsmtUnfSF: Unfinished square feet of basement area
- TotalBsmtSF: Total square feet of basement area
- Heating: Type of heating
- HeatingQC: Heating quality and condition

- CentralAir: Central air conditioning
- Electrical: Electrical system
- 1stFlrSF: First Floor square feet
- 2ndFlrSF: Second floor square feet
- LowQualFinSF: Low quality finished square feet (all floors)
- GrLivArea: Above grade (ground) living area square feet
- BsmtFullBath: Basement full bathrooms
- BsmtHalfBath: Basement half bathrooms
- FullBath: Full bathrooms above grade
- HalfBath: Half baths above grade
- Bedroom: Number of bedrooms above basement level
- Kitchen: Number of kitchens
- KitchenQual: Kitchen quality
- TotRmsAbvGrd: Total rooms above grade (does not include bathrooms)
- Functional: Home functionality rating
- Fireplaces: Number of fireplaces
- FireplaceQu: Fireplace quality
- -GarageType: Garage location
- GarageYrBlt: Year garage was built
- GarageFinish: Interior finish of the garage
- GarageCars: Size of garage in car capacity
- GarageArea: Size of garage in square feet
- GarageQual: Garage quality
- GarageCond: Garage condition
- PavedDrive: Paved driveway
- WoodDeckSF: Wood deck area in square feet
- OpenPorchSF: Open porch area in square feet
- EnclosedPorch: Enclosed porch area in square feet
- 3SsnPorch: Three season porch area in square feet
- ScreenPorch: Screen porch area in square feet
- PoolArea: Pool area in square feet
- PoolQC: Pool quality
- Fence: Fence quality
- MiscFeature: Miscellaneous feature not covered in other categories
- MiscVal: \$Value of miscellaneous feature
- MoSold: Month Sold
- YrSold: Year Sold
- SaleType: Type of sale
- SaleCondition: Condition of sale

```
In [15]: def data_process(df):
    all_df = pd.DataFrame(index = df.index)

    all_df["LotFrontage"] = df["LotFrontage"]
    for key, group in lot_frontage_by_neighborhood:
        #Filling in missing LotFrontage values by the median
        idx = (df["Neighborhood"] == key) & (df["LotFrontage"].isnull())
        all_df.loc[idx, "LotFrontage"] = group.median()
        all_df["LotArea"] = df["LotArea"]

    all_df["MasVnrArea"] = df["MasVnrArea"]
    all_df["MasVnrArea"].fillna(0, inplace=True)

    all_df["BsmtFinSF1"] = df["BsmtFinSF1"]
    all_df["BsmtFinSF1"].fillna(0, inplace=True)

    all_df["BsmtFinSF2"] = df["BsmtFinSF2"]
    all_df["BsmtFinSF2"].fillna(0, inplace=True)

    all_df["BsmtUnfSF"] = df["BsmtUnfSF"]
    all_df["BsmtUnfSF"].fillna(0, inplace=True)

    all_df["TotalBsmtSF"] = df["TotalBsmtSF"]
    all_df["TotalBsmtSF"].fillna(0, inplace=True)

    all_df["1stFlrSF"] = df["1stFlrSF"]
    all_df["2ndFlrSF"] = df["2ndFlrSF"]
    all_df["GrLivArea"] = df["GrLivArea"]

    all_df["GarageArea"] = df["GarageArea"]
    all_df["GarageArea"].fillna(0, inplace=True)

    all_df["WoodDeckSF"] = df["WoodDeckSF"]
    all_df["OpenPorchSF"] = df["OpenPorchSF"]
    all_df["EnclosedPorch"] = df["EnclosedPorch"]
    all_df["3SsnPorch"] = df["3SsnPorch"]
    all_df["ScreenPorch"] = df["ScreenPorch"]

    all_df["BsmtFullBath"] = df["BsmtFullBath"]
    all_df["BsmtFullBath"].fillna(0, inplace=True)

    all_df["BsmtHalfBath"] = df["BsmtHalfBath"]
    all_df["BsmtHalfBath"].fillna(0, inplace=True)

    all_df["FullBath"] = df["FullBath"]
    all_df["HalfBath"] = df["HalfBath"]
    all_df["BedroomAbvGr"] = df["BedroomAbvGr"]
    all_df["KitchenAbvGr"] = df["KitchenAbvGr"]
    all_df["TotRmsAbvGrd"] = df["TotRmsAbvGrd"]
    all_df["Fireplaces"] = df["Fireplaces"]
```

```

all_df["GarageCars"] = df["GarageCars"]
all_df["GarageCars"].fillna(0, inplace=True)

all_df["CentralAir"] = (df["CentralAir"] == "Y") * 1.0

all_df["OverallQual"] = df["OverallQual"]
all_df["OverallCond"] = df["OverallCond"]

"""following case are ordinal so we are performing label encoding here"""

nan = float('nan')
qual_dict = {nan: 0, "NA": 0, "Po": 1, "Fa": 2, "TA": 3, "Gd": 4, "Ex": 5}
all_df["ExterQual"] = df["ExterQual"].map(qual_dict).astype(int)
all_df["ExterCond"] = df["ExterCond"].map(qual_dict).astype(int)
all_df["BsmtQual"] = df["BsmtQual"].map(qual_dict).astype(int)
all_df["BsmtCond"] = df["BsmtCond"].map(qual_dict).astype(int)
all_df["HeatingQC"] = df["HeatingQC"].map(qual_dict).astype(int)
all_df["KitchenQual"] = df["KitchenQual"].map(qual_dict).astype(int)
all_df["FireplaceQu"] = df["FireplaceQu"].map(qual_dict).astype(int)
all_df["GarageQual"] = df["GarageQual"].map(qual_dict).astype(int)
all_df["GarageCond"] = df["GarageCond"].map(qual_dict).astype(int)

all_df["BsmtExposure"] = df["BsmtExposure"].map(
    {nan: 0, "No": 1, "Mn": 2, "Av": 3, "Gd": 4}).astype(int)

bsmt_fin_dict = {nan: 0, "Unf": 1, "LwQ": 2, "Rec": 3, "BLQ": 4, "ALQ": 5, "GLQ": 6}
all_df["BsmtFinType1"] = df["BsmtFinType1"].map(bsmt_fin_dict).astype(int)
all_df["BsmtFinType2"] = df["BsmtFinType2"].map(bsmt_fin_dict).astype(int)

all_df["Functional"] = df["Functional"].map(
    {nan: 0, "Sal": 1, "Sev": 2, "Maj2": 3, "Maj1": 4,
     "Mod": 5, "Min2": 6, "Min1": 7, "Typ": 8}).astype(int)

all_df["GarageFinish"] = df["GarageFinish"].map(
    {nan: 0, "Unf": 1, "RFn": 2, "Fin": 3}).astype(int)

all_df["Fence"] = df["Fence"].map(
    {nan: 0, "MnWw": 1, "GdWo": 2, "MnPrv": 3, "GdPrv": 4}).astype(int)

all_df["PoolQC"] = df["PoolQC"].map(qual_dict).astype(int)

all_df["YearBuilt"] = df["YearBuilt"]
all_df["YearRemodAdd"] = df["YearRemodAdd"]

all_df["GarageYrBlt"] = df["GarageYrBlt"]
all_df["GarageYrBlt"].fillna(0.0, inplace=True)

all_df["MoSold"] = df["MoSold"]
all_df["YrSold"] = df["YrSold"]

```

```

all_df["LowQualFinSF"] = df["LowQualFinSF"]
all_df["MiscVal"] = df["MiscVal"]

all_df["PoolQC"] = df["PoolQC"].map(qual_dict).astype(int)

all_df["PoolArea"] = df["PoolArea"]
all_df["PoolArea"].fillna(0, inplace=True)

# Add categorical features as numbers too. It seems to help a bit.
all_df = factorize(df, all_df, "MSSubClass")
all_df = factorize(df, all_df, "MSZoning", "RL")
all_df = factorize(df, all_df, "LotConfig")
all_df = factorize(df, all_df, "Neighborhood")
all_df = factorize(df, all_df, "Condition1")
all_df = factorize(df, all_df, "BldgType")
all_df = factorize(df, all_df, "HouseStyle")
all_df = factorize(df, all_df, "RoofStyle")
all_df = factorize(df, all_df, "Exterior1st", "Other")
all_df = factorize(df, all_df, "Exterior2nd", "Other")
all_df = factorize(df, all_df, "MasVnrType", "None")
all_df = factorize(df, all_df, "Foundation")
all_df = factorize(df, all_df, "SaleType", "Oth")
all_df = factorize(df, all_df, "SaleCondition")

"""In following code I am converting values of those features as 0 or 1"""

# IR2 and IR3 don't appear that often, so just make a distinction
# between regular and irregular.
all_df["IsRegularLotShape"] = (df["LotShape"] == "Reg") * 1

# Most properties are level; bin the other possibilities together
# as "not level".
all_df["IsLandLevel"] = (df["LandContour"] == "Lvl") * 1

# Most land slopes are gentle; treat the others as "not gentle".
all_df["IsLandSlopeGentle"] = (df["LandSlope"] == "Gtl") * 1

# Most properties use standard circuit breakers.
all_df["IsElectricalSBrkr"] = (df["Electrical"] == "SBrkr") * 1

# About 2/3rd have an attached garage.
all_df["IsGarageDetached"] = (df["GarageType"] == "Detchd") * 1

# Most have a paved drive. Treat dirt/gravel and partial pavement
# as "not paved".
all_df["IsPavedDrive"] = (df["PavedDrive"] == "Y") * 1

# The only interesting "misc. feature" is the presence of a shed.
all_df["HasShed"] = (df["MiscFeature"] == "Shed") * 1.

# If YearRemodAdd != YearBuilt, then a remodeling took place at some point.

```

```

all_df["Remodeled"] = (all_df["YearRemodAdd"] != all_df["YearBuilt"]) * 1

# Did a remodeling happen in the year the house was sold?
all_df["RecentRemodel"] = (all_df["YearRemodAdd"] == all_df["YrSold"]) * 1

# Was this house sold in the year it was built?
all_df["VeryNewHouse"] = (all_df["YearBuilt"] == all_df["YrSold"]) * 1

all_df["Has2ndFloor"] = (all_df["2ndFlrSF"] == 0) * 1
all_df["HasMasVnr"] = (all_df["MasVnrArea"] == 0) * 1
all_df["HasWoodDeck"] = (all_df["WoodDeckSF"] == 0) * 1
all_df["HasOpenPorch"] = (all_df["OpenPorchSF"] == 0) * 1
all_df["HasEnclosedPorch"] = (all_df["EnclosedPorch"] == 0) * 1
all_df["Has3SsnPorch"] = (all_df["3SsnPorch"] == 0) * 1
all_df["HasScreenPorch"] = (all_df["ScreenPorch"] == 0) * 1

# Months with the largest number of deals may be significant.
# mx = max(train["MoSold"].groupby(train["MoSold"]).count())
# all_df["HighSeason"] = df["MoSold"].replace(
#     train["MoSold"].groupby(train["MoSold"]).count()/mx)

# mx = max(train["MSSubClass"].groupby(train["MSSubClass"]).count())
# all_df["NewerDwelling"] = df["MSSubClass"].replace(
#     train["MSSubClass"].groupby(train["MSSubClass"]).count()/mx)

# following portion was calculated with above commented part of the code.
# Instead of the fraction value putting binary value helps for generalization
all_df["HighSeason"] = df["MoSold"].replace(
    {1: 0, 2: 0, 3: 0, 4: 1, 5: 1, 6: 1, 7: 1, 8: 0, 9: 0, 10: 0, 11: 0, 12: 0})

all_df["NewerDwelling"] = df["MSSubClass"].replace(
    {20: 1, 30: 0, 40: 0, 45: 0, 50: 0, 60: 1, 70: 0, 75: 0, 80: 0, 85: 0,
     90: 0, 120: 1, 150: 0, 160: 0, 180: 0, 190: 0})

all_df.loc[df.Neighborhood == 'NridgHt', "Neighborhood_Good"] = 1
all_df.loc[df.Neighborhood == 'Crawfor', "Neighborhood_Good"] = 1
all_df.loc[df.Neighborhood == 'StoneBr', "Neighborhood_Good"] = 1
all_df.loc[df.Neighborhood == 'Somerst', "Neighborhood_Good"] = 1
all_df.loc[df.Neighborhood == 'NoRidge', "Neighborhood_Good"] = 1
all_df["Neighborhood_Good"].fillna(0, inplace=True)

# House completed before sale or not
all_df["SaleCondition_PriceDown"] = df.SaleCondition.replace(
    {'Abnorml': 1, 'Alloca': 1, 'AdjLand': 1, 'Family': 1, 'Normal': 0, 'Partial': 0})

# House completed before sale or not
all_df["BoughtOffPlan"] = df.SaleCondition.replace(
    {"Abnorml" : 0, "Alloca" : 0, "AdjLand" : 0, "Family" : 0, "Normal" : 0, "Partial" : 1})

all_df["BadHeating"] = df.HeatingQC.replace(

```

```

{'Ex': 0, 'Gd': 0, 'TA': 0, 'Fa': 1, 'Po': 1})

area_cols = ['LotFrontage', 'LotArea', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
             'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'GrLivArea', 'GarageArea', 'WoodDeckSF',
             'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'LowQualFinSF', 'PoolArea' ]
all_df["TotalArea"] = all_df[area_cols].sum(axis=1)

all_df["TotalArea1st2nd"] = all_df["1stFlrSF"] + all_df["2ndFlrSF"]

all_df["Age"] = 2010 - all_df["YearBuilt"]
all_df["TimeSinceSold"] = 2010 - all_df["YrSold"]

all_df["SeasonSold"] = all_df["MoSold"].map({12:0, 1:0, 2:0, 3:1, 4:1, 5:1,
                                             6:2, 7:2, 8:2, 9:3, 10:3, 11:3}).astype(int)

all_df["YearsSinceRemodel"] = all_df["YrSold"] - all_df["YearRemodAdd"]

# Simplifications of existing features into bad/average/good.
all_df["SimplOverallQual"] = all_df.OverallQual.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2, 6 : 2, 7 : 3, 8 : 3, 9 : 3, 10 : 3})
all_df["SimplOverallCond"] = all_df.OverallCond.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2, 6 : 2, 7 : 3, 8 : 3, 9 : 3, 10 : 3})
all_df["SimplPoolQC"] = all_df.PoolQC.replace(
    {1 : 1, 2 : 1, 3 : 2, 4 : 2})
all_df["SimplGarageCond"] = all_df.GarageCond.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})
all_df["SimplGarageQual"] = all_df.GarageQual.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})
all_df["SimplFireplaceQu"] = all_df.FireplaceQu.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})
all_df["SimplFireplaceQu"] = all_df.FireplaceQu.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})
all_df["SimplFunctional"] = all_df.Functional.replace(
    {1 : 1, 2 : 1, 3 : 2, 4 : 2, 5 : 3, 6 : 3, 7 : 3, 8 : 4})
all_df["SimplKitchenQual"] = all_df.KitchenQual.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})
all_df["SimplHeatingQC"] = all_df.HeatingQC.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})
all_df["SimplBsmtFinType1"] = all_df.BsmtFinType1.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2, 6 : 2})
all_df["SimplBsmtFinType2"] = all_df.BsmtFinType2.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2, 6 : 2})
all_df["SimplBsmtCond"] = all_df.BsmtCond.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})
all_df["SimplBsmtQual"] = all_df.BsmtQual.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})
all_df["SimplExterCond"] = all_df.ExterCond.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})
all_df["SimplExterQual"] = all_df.ExterQual.replace(
    {1 : 1, 2 : 1, 3 : 1, 4 : 2, 5 : 2})

```

```
# Bin by neighborhood (a little arbitrarily). Values were computed by:
# train_df["SalePrice"].groupby(train_df["Neighborhood"]).median().sort_values()
neighborhood_map = {
    "MeadowV" : 0, # 88000
    "IDOTRR" : 1, # 103000
    "BrDale" : 1, # 106000
    "OldTown" : 1, # 119000
    "Edwards" : 1, # 119500
    "BrkSide" : 1, # 124300
    "Sawyer" : 1, # 135000
    "Blueste" : 1, # 137500
    "SWISU" : 2, # 139500
    "NAmes" : 2, # 140000
    "NPkVill" : 2, # 146000
    "Mitchel" : 2, # 153500
    "SawyerW" : 2, # 179900
    "Gilbert" : 2, # 181000
    "NWAmes" : 2, # 182900
    "Blmngtn" : 2, # 191000
    "CollgCr" : 2, # 197200
    "ClearCr" : 3, # 200250
    "Crawfor" : 3, # 200624
    "Veenker" : 3, # 218000
    "Somerst" : 3, # 225500
    "Timber" : 3, # 228475
    "StoneBr" : 4, # 278000
    "NoRidge" : 4, # 290000
    "NridgHt" : 4, # 315000
}

all_df["NeighborhoodBin"] = df["Neighborhood"].map(neighborhood_map)
return all_df
```

In the above block I have done following operations :

- Filled with 0 for some features like MasVnrArea , GarageArea , TotalBsmtSF etc .According to the documentation of the dataset if these features have any field empty then that means the feature is not available. So I have done this operation according to documentation of the dataset.
- CentralAir feature was given has two field only 'Y' or 'N' so I have converted that to 0 or 1
- For some ordinal features I ave performed lable encoding. Those features are ExterQual , ExterCond, BsmtQual, BsmtCond, HeatingQC, KitchenQual etc
- I have converted some features from categorical to numerical and those features are MSSubClass, MSZoning , LotConfig, RL , LotConfig,Neighborhood, Condition1 ,BldgType, HouseStyle , HouseStyle, Exterior1st, Other, Exterior2nd, MasVnrType, Foundation, SaleType and SaleCondition
- Converted fields of some Features to 0 or 1 based on the understanding of the dataset and a little bit research. Those features with the changing reasons are given below
 - IsRegularLotShape : Field IR2 and IR3 don't appear that often, so just make a distinction between regular and irregular.
 - IsLandLevel : Most land slopes are gentle; treat the others as "not gentle".
 - IsElectricalSBrkr : Most properties use standard circuit breakers.
 - IsGarageDetached : About 2/3rd have an attached garage.
 - IsPavedDrive : Most have a paved drive. Treat dirt/gravel and partial pavement as "not paved".
 - HasShed : The only interesting "misc. feature" is the presence of a shed.
 - Remodeled : If YearRemodAdd != YearBuilt, then a remodeling took place at some point.
 - RecentRemodel : Did a remodeling happen in the year the house was sold?
 - VeryNewHouse : Was this house sold in the year it was built?
 - sofe other features dont need to describe they are self explanatory Has2ndFloor , HasMasVnr , HasWoodDeck ,HasOpenPorch ,HasEnclosedPorch ,Has3SsnPorch , HasScreenPorch
- Simplifications of existing features into bad/average/good. Features : SimplOverallQual, SimplOverallCond, SimplPoolQC ,SimplGarageCond, SimplGarageQual, SimplFireplaceQu ,SimplFunctional ,SimplKitchenQual, SimplHeatingQC ,SimplBsmtFinType1, SimplBsmtFinType2 ,SimplBsmtCond , SimplBsmtQual ,SimplExterCond ,SimplExterQual.
- mapped neighborhood based on their quality.The mapping is as followed:


```
"MeadowV" : 0, # 88000
"IDOTRR" : 1, # 103000
"BrDale" : 1, # 106000
"OldTown" : 1, # 119000
"Edwards" : 1, # 119500
"BrkSide" : 1, # 124300
"Sawyer" : 1, # 135000
"Riviera" : 1, # 137500
```

```
In [16]: train_processed = data_process(train)
         test_processed = data_process(test)

         print("shape of train :", train_processed.shape)
         print("shape of test :", test_processed.shape)

         shape of train : (1458, 111)
         shape of test : (1459, 111)
```

Keeping NeighborhoodBin into a temporary DataFrame because we want to use the unscaled version later on (to one-hot encode it).

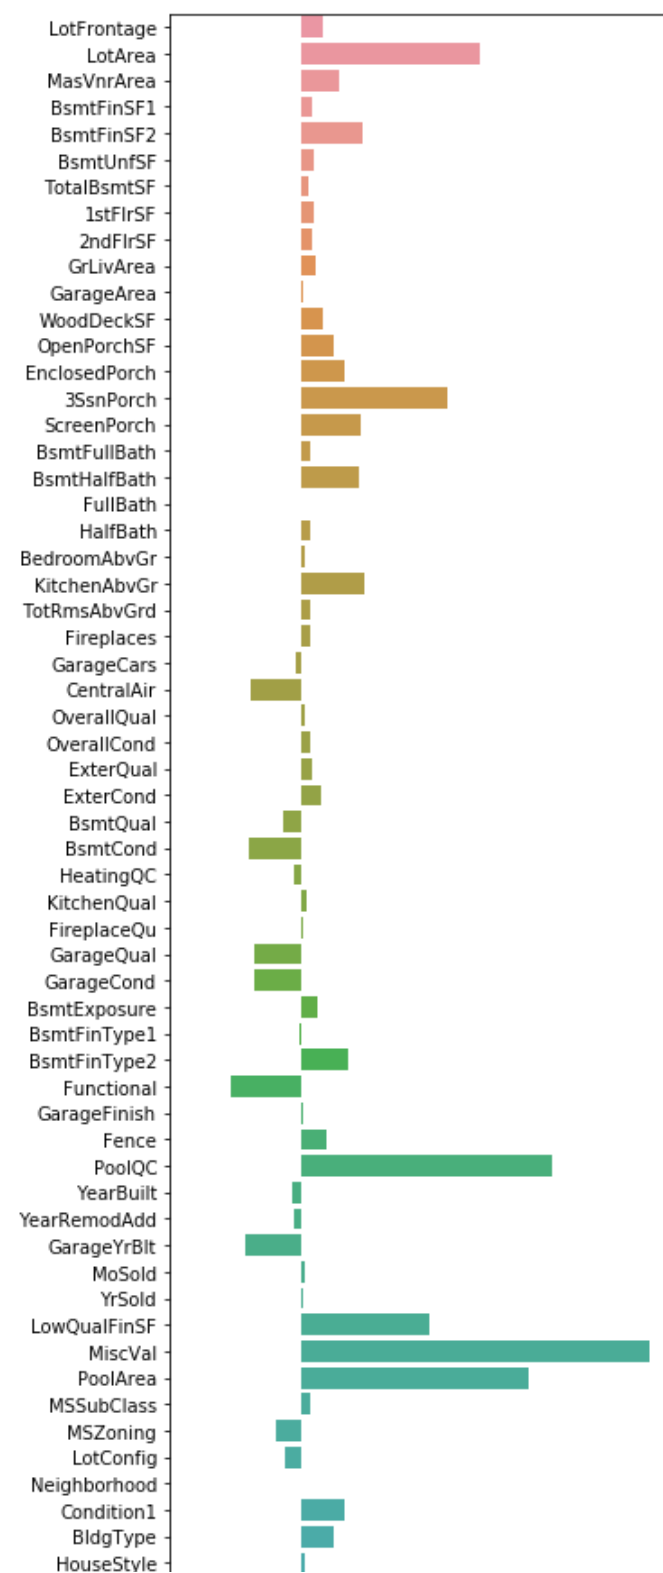
```
In [17]: # Keeping NeighborhoodBin into a temporary DataFrame because we want to use the
         # unscaled version later on (to one-hot encode it).
         neighborhood_bin_train = pd.DataFrame(index = train.index)
         neighborhood_bin_train["NeighborhoodBin"] = train_processed["NeighborhoodBin"]
         neighborhood_bin_test = pd.DataFrame(index = test.index)
         neighborhood_bin_test["NeighborhoodBin"] = test_processed["NeighborhoodBin"]
```

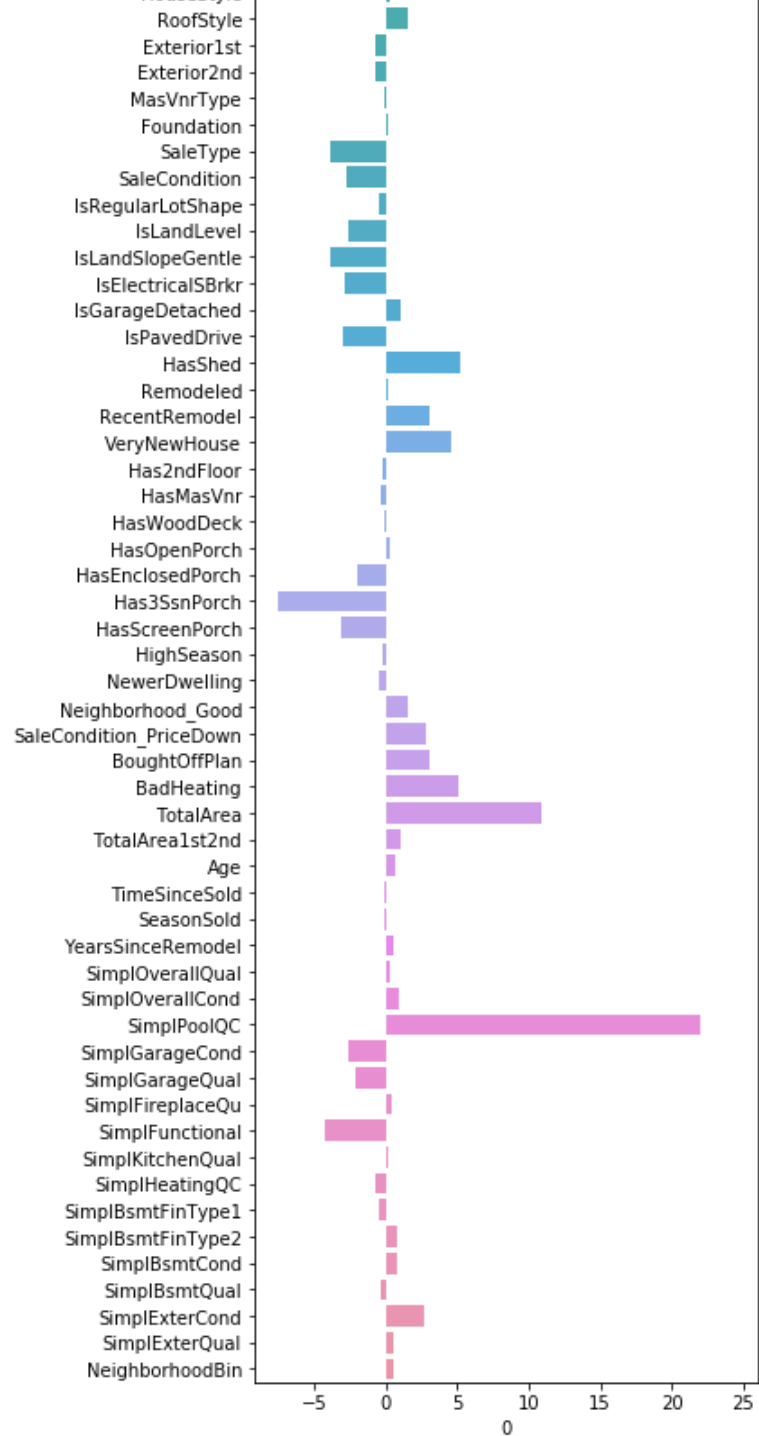
Skewness & Normalization

skewness train set

```
In [18]: from scipy.stats import skew
import seaborn as sns
numeric_features = train_processed.dtypes[train_processed.dtypes != "object"].index

skewness = train_processed[numeric_features].skew(axis=0 , skipna =True)
skewness = pd.DataFrame(skewness)
plt.figure(figsize=[5,30])
# skw = sns.load_dataset(skewness)
ax = sns.barplot( y= skewness.index , x=skewness[0] , data = skewness)
plt.show()
# print('skew: ',train_processed[numeric_features].skew())
```





```
In [19]: numeric_features = train_processed.dtypes[train_processed.dtypes != "object"].index
```

```
# Transform the skewed numeric features by taking log(feature + 1).  
# This will make the features more normal.  
from scipy.stats import skew  
  
skewed = train_processed[numeric_features].apply(lambda x: skew(x.dropna().astype(float)))  
skewed = skewed[(skewed < -0.75) | (skewed > 0.75)]  
skewed = skewed.index  
  
train_processed[skewed] = np.log1p(train_processed[skewed])  
  
# Additional processing: scale the data.  
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
scaled = scaler.fit_transform(train_processed[numeric_features])  
  
for i, col in enumerate(numeric_features):  
    train_processed[col] = scaled[:, i]
```


```
/home/navid/anaconda3/envs/tf/lib/python3.6/site-packages/sklearn/preprocessing/data.py:645: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by StandardScaler.
```

```
    return self.partial_fit(X, y)
```

```
/home/navid/anaconda3/envs/tf/lib/python3.6/site-packages/sklearn/base.py:464: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by StandardScaler.
```

```
    return self.fit(X, **fit_params).transform(X)
```

```
In [20]: from scipy.stats import skew
numeric_features = train_processed.dtypes[train_processed.dtypes != "object"].index
skewness = train_processed[numeric_features].skew(axis=0 , skipna =True)
skewness = pd.DataFrame(skewness)
plt.figure(figsize=[5,30])
# skw = sns.load_dataset(skewness)
ax = sns.barplot( y= skewness.index , x=skewness[0] , data = skewness)
plt.show()
```



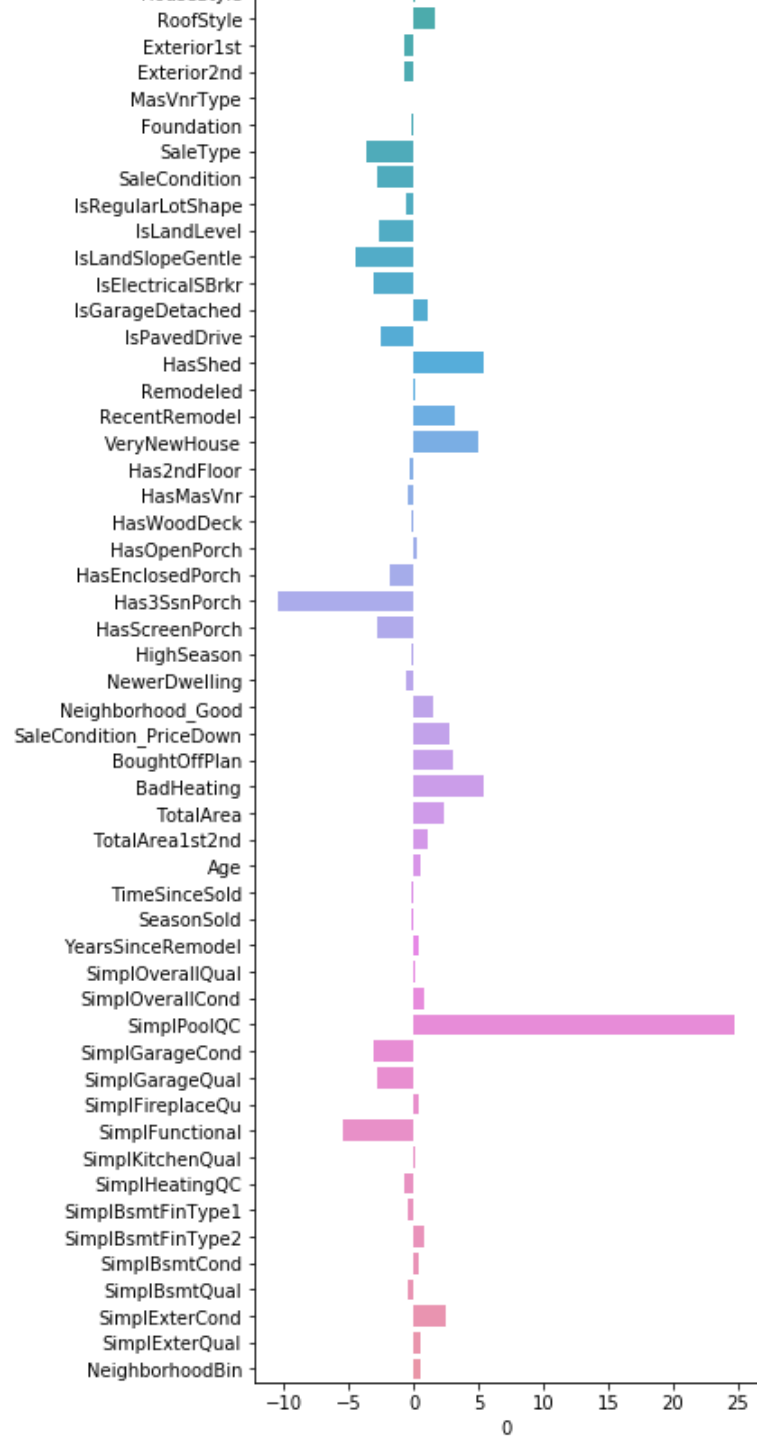


Test Skewness

```
In [21]: numeric_features = test_processed.dtypes[train_processed.dtypes != "object"].index
skewness = test_processed[numeric_features].skew(axis=0 , skipna =True)
skewness = pd.DataFrame(skewness)

plt.figure(figsize=[5,30])
# skw = sns.load_dataset(skewness)
ax = sns.barplot( y= skewness.index , x=skewness[0] , data = skewness)
plt.show()
# print('skew: ',test_processed[numeric_features].skew())
```





```
In [22]: numeric_features = test_processed.dtypes[train_processed.dtypes != "object"].index
```

```
# Transform the skewed numeric features by taking log(feature + 1).  
# This will make the features more normal.  
from scipy.stats import skew  
  
skewed = test_processed[numeric_features].apply(lambda x: skew(x.dropna().astype(float)))  
skewed = skewed[(skewed < -0.75) | (skewed > 0.75)]  
skewed = skewed.index  
  
test_processed[skewed] = np.log1p(test_processed[skewed])  
  
# Additional processing: scale the data.  
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
  
scaled = scaler.fit_transform(test_processed[numeric_features])  
for i, col in enumerate(numeric_features):  
    test_processed[col] = scaled[:, i]
```

```
/home/navid/anaconda3/envs/tf/lib/python3.6/site-packages/sklearn/preprocessing/data.py:645: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by StandardScaler.
```

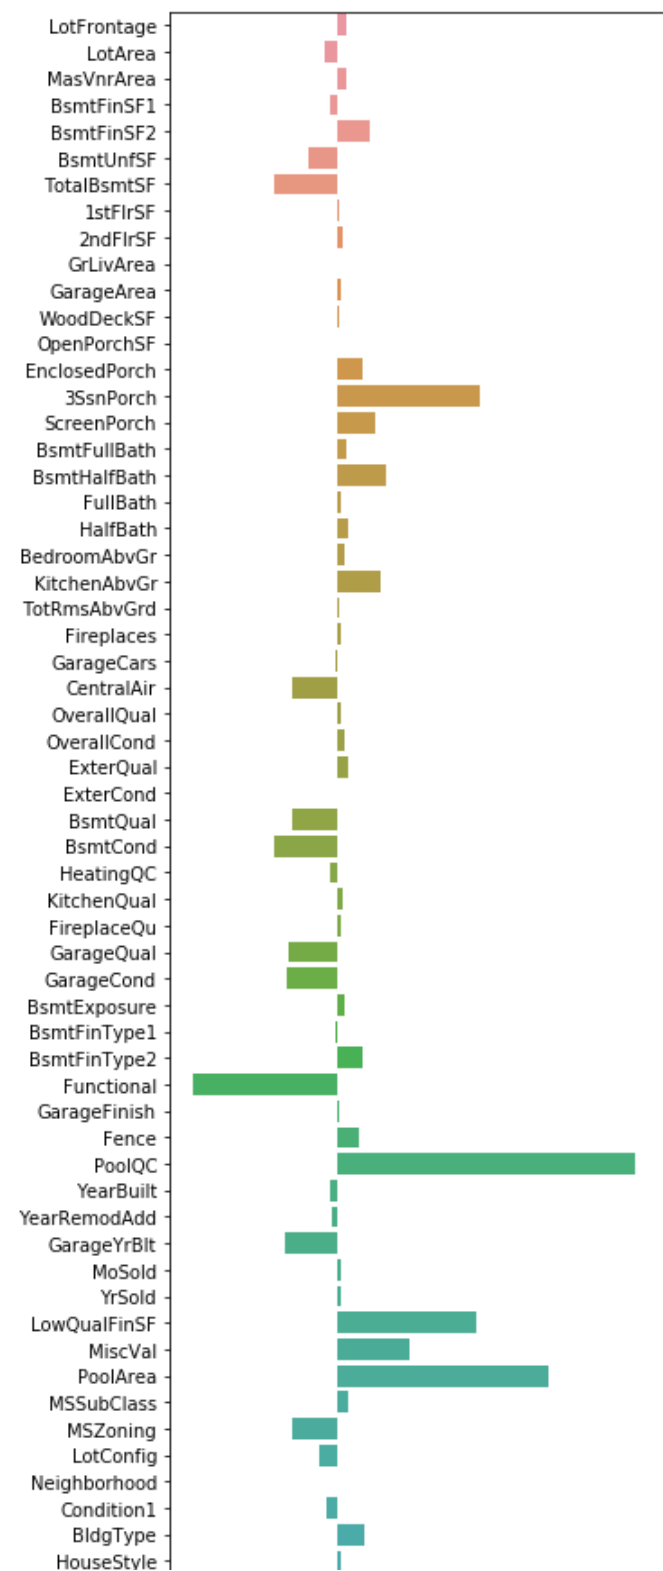
```
    return self.partial_fit(X, y)
```

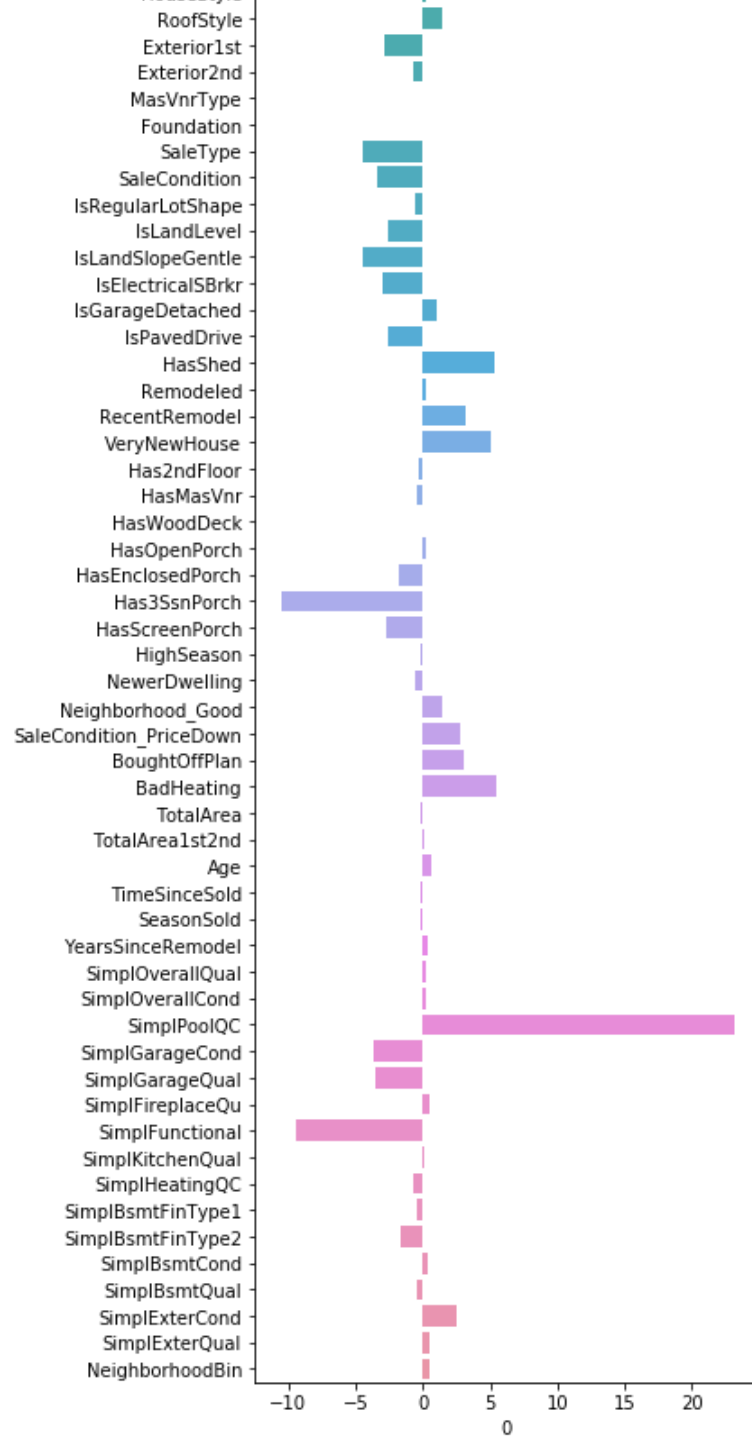
```
/home/navid/anaconda3/envs/tf/lib/python3.6/site-packages/sklearn/base.py:464: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by StandardScaler.
```

```
    return self.fit(X, **fit_params).transform(X)
```

```
In [23]: numeric_features = test_processed.dtypes[train_processed.dtypes != "object"].index
skewness = test_processed[numeric_features].skew(axis=0 , skipna =True)
skewness = pd.DataFrame(skewness)

plt.figure(figsize=[5,30])
# skw = sns.load_dataset(skewness)
ax = sns.barplot( y= skewness.index , x=skewness[0] , data = skewness)
plt.show()
# print('skew: ',test_processed[numeric_features].skew())
```





Additional processing to scale the data.

One hot encoding

To encode categorical integer features as a one-hot numeric array we are using one hot encoding. This will transform each value of categories into a features and make those a column value of dataframe. Finally put binary values in the rows of those column.

In [24]: # for example:

```
#
```

CompanyName	Categoricalvalue	Price
VW	1	20000
Acura	2	10011
Honda	3	50000
Honda	3	10000

```
#
```

converting it to one Hot encoding:

```
#
```

VW	Acura	Honda	Price
1	0	0	20000
0	1	0	10011
0	0	1	50000
0	0	1	10000

```
#
```

reference: <https://hackernoon.com/what-is-one-hot-encoding-why-and-when-do-you-have-to-use-it-e3c6186d008f>

In this section at first we merge train and test data (variable name predictor_cols and predictor_cols_test). We did it because there is some features in train data which is missing in test data again same thing can happen for test data too.

```
In [25]: # Convert categorical features using one-hot encoding.
def onehot(onehot_df, df, column_name, fill_na, drop_name):
    onehot_df[column_name] = df[column_name]
    if fill_na is not None:
        onehot_df[column_name].fillna(fill_na, inplace=True)

    dummies = pd.get_dummies(onehot_df[column_name], prefix="_" + column_name)

    onehot_df = onehot_df.join(dummies)
    onehot_df = onehot_df.drop([column_name], axis=1)
    return onehot_df
```

performing one hot

```

In [26]: def proceed_onehot(df):
onehot_df = pd.DataFrame(index = df.index)

onehot_df = onehot(onehot_df, df, "MSSubClass", None, "40")
onehot_df = onehot(onehot_df, df, "MSZoning", "RL", "RH")
onehot_df = onehot(onehot_df, df, "LotConfig", None, "FR3")
onehot_df = onehot(onehot_df, df, "Neighborhood", None, "OldTown")
onehot_df = onehot(onehot_df, df, "Condition1", None, "RRNe")
onehot_df = onehot(onehot_df, df, "BldgType", None, "2fmCon")
onehot_df = onehot(onehot_df, df, "HouseStyle", None, "1.5Unf")
onehot_df = onehot(onehot_df, df, "RoofStyle", None, "Shed")
onehot_df = onehot(onehot_df, df, "Exterior1st", "VinylSd", "CBlock")
onehot_df = onehot(onehot_df, df, "Exterior2nd", "VinylSd", "CBlock")
onehot_df = onehot(onehot_df, df, "Foundation", None, "Wood")
onehot_df = onehot(onehot_df, df, "SaleType", "WD", "Oth")
onehot_df = onehot(onehot_df, df, "SaleCondition", "Normal", "AdjLand")

# Fill in missing MasVnrType for rows that do have a MasVnrArea.
temp_df = df[["MasVnrType", "MasVnrArea"]].copy()
idx = (df["MasVnrArea"] != 0) & ((df["MasVnrType"] == "None") | (df["MasVnrType"].isnull()))
temp_df.loc[idx, "MasVnrType"] = "BrkFace"
onehot_df = onehot(onehot_df, temp_df, "MasVnrType", "None", "BrkCmn")

# Also add the booleans from calc_df as dummy variables.
onehot_df = onehot(onehot_df, df, "LotShape", None, "IR3")
onehot_df = onehot(onehot_df, df, "LandContour", None, "Low")
onehot_df = onehot(onehot_df, df, "LandSlope", None, "Sev")
onehot_df = onehot(onehot_df, df, "Electrical", "SBrkr", "FuseP")
onehot_df = onehot(onehot_df, df, "GarageType", "None", "CarPort")
onehot_df = onehot(onehot_df, df, "PavedDrive", None, "P")
onehot_df = onehot(onehot_df, df, "MiscFeature", "None", "Othr")

# Features we can probably ignore (but want to include anyway to see
# if they make any positive difference).
# Definitely ignoring Utilities: all records are "AllPub", except for
# one "NoSeWa" in the train set and 2 NA in the test set.
onehot_df = onehot(onehot_df, df, "Street", None, "Grvl")
onehot_df = onehot(onehot_df, df, "Alley", "None", "Grvl")
onehot_df = onehot(onehot_df, df, "Condition2", None, "PosA")
onehot_df = onehot(onehot_df, df, "RoofMatl", None, "WdShake")
onehot_df = onehot(onehot_df, df, "Heating", None, "Wall")

# I have these as numerical variables too.
onehot_df = onehot(onehot_df, df, "ExterQual", "None", "Ex")
onehot_df = onehot(onehot_df, df, "ExterCond", "None", "Ex")
onehot_df = onehot(onehot_df, df, "BsmtQual", "None", "Ex")
onehot_df = onehot(onehot_df, df, "BsmtCond", "None", "Ex")
onehot_df = onehot(onehot_df, df, "HeatingQC", "None", "Ex")
onehot_df = onehot(onehot_df, df, "KitchenQual", "TA", "Ex")
onehot_df = onehot(onehot_df, df, "FireplaceQu", "None", "Ex")
onehot_df = onehot(onehot_df, df, "GarageQual", "None", "Ex")

```

```

onehot_df = onehot(onehot_df, df, "GarageCond", "None", "Ex")
onehot_df = onehot(onehot_df, df, "PoolQC", "None", "Ex")
onehot_df = onehot(onehot_df, df, "BsmtExposure", "None", "Gd")
onehot_df = onehot(onehot_df, df, "BsmtFinType1", "None", "GLQ")
onehot_df = onehot(onehot_df, df, "BsmtFinType2", "None", "GLQ")
onehot_df = onehot(onehot_df, df, "Functional", "Typ", "Typ")
onehot_df = onehot(onehot_df, df, "GarageFinish", "None", "Fin")
onehot_df = onehot(onehot_df, df, "Fence", "None", "MnPrv")
onehot_df = onehot(onehot_df, df, "MoSold", None, None)

# Divide up the years between 1871 and 2010 in slices of 20 years.
year_map = pd.concat(pd.Series("YearBin" + str(i+1), index=range(1871+i*20,1891+i*20)) for i in range(0, 7))

yearbin_df = pd.DataFrame(index = df.index)
yearbin_df["GarageYrBltBin"] = df.GarageYrBlt.map(year_map)
yearbin_df["GarageYrBltBin"].fillna("NoGarage", inplace=True)

yearbin_df["YearBuiltBin"] = df.YearBuilt.map(year_map)
yearbin_df["YearRemodAddBin"] = df.YearRemodAdd.map(year_map)

onehot_df = onehot(onehot_df, yearbin_df, "GarageYrBltBin", None, None)
onehot_df = onehot(onehot_df, yearbin_df, "YearBuiltBin", None, None)
onehot_df = onehot(onehot_df, yearbin_df, "YearRemodAddBin", None, None)

return onehot_df

# Add the one-hot encoded categorical features.
onehot_df = proceed_onehot(train)
onehot_df = onehot(onehot_df, neighborhood_bin_train, "NeighborhoodBin", None, None)
train_processed = train_processed.join(onehot_df)

```

These onehot columns are missing in the test data, so drop them from the training data or we might overfit on them.

```

In [27]: drop_cols = [
    "_Exterior1st_ImStucc", "_Exterior1st_Stone",
    "_Exterior2nd_Other", "_HouseStyle_2.5Fin",

    "_RoofMatl_Membran", "_RoofMatl_Metal", "_RoofMatl_Roll",
    "_Condition2_RRAe", "_Condition2_RRAn", "_Condition2_RRNn",
    "_Heating_Floor", "_Heating_OthW",

    "_Electrical_Mix",
    "_MiscFeature_TenC",
    "_GarageQual_Ex", "_PoolQC_Fa"
]
train_processed.drop(drop_cols, axis=1, inplace=True)

```

```
In [28]: onehot_df = proceed_onehot(test)
onehot_df = onehot(onehot_df, neighborhood_bin_test, "NeighborhoodBin", None, None)
test_processed = test_processed.join(onehot_df)
```

This column is missing in the training data. There is only one example with this value in the test set. So just drop it.

```
In [29]: test_processed.drop(["_MSSubClass_150"], axis=1, inplace=True)
```

Drop these columns. They are either not very helpful or they cause overfitting.

```
In [30]: drop_cols = [
    "_Condition2_PosN",      # only two are not zero
    "_MSZoning_C (all)",
    "_MSSubClass_160",
]
train_processed.drop(drop_cols, axis=1, inplace=True)
test_processed.drop(drop_cols, axis=1, inplace=True)
```

log transform

We take the log here because the error metric is between the log of the SalePrice and the log of the predicted price. That does mean we need to `exp()` the prediction to get an actual sale price.

```
In [31]: target = pd.DataFrame(index = train_processed.index, columns=["SalePrice"])
target["SalePrice"] = np.log(train["SalePrice"])
# train_processed.drop(["SalePrice"], axis=1, inplace=True)

print("Training set size:", train_processed.shape)
print("Test set size:", test_processed.shape)
```

```
Training set size: (1458, 403)
Test set size: (1459, 403)
```

Split Data for training and testing

```
In [32]: X_train, X_test, y_train, y_test = train_test_split(train_processed,
    target,
    # train_size = 0.99,
    test_size = 0.2,
    random_state = 0)
```

```
In [33]: prediction_dict = dict()
submit_prediction_dict = dict()

submit = False
save_score = False

if submit :
    X_train = train_processed
    y_train = target
else:
    X_train = X_train
    y_train = y_train
```

Testing different models

Random Forest Regressor

A random forest is a meta estimator that fits a number of classifying decision trees on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting.

```
In [34]: my_model = RandomForestRegressor(n_estimators=500,n_jobs=-1)

my_model.fit(X_train, y_train)
prediction = my_model.predict(X_test)
if submit:
    submit_prediction = my_model.predict(test_processed)
    submit_prediction_dict['Random Forest Regressor'] = submit_prediction

prediction_dict['Random Forest Regressor'] = prediction

print('root mean absolute error: ',rmse(y_test, prediction))
print('accuracy score: ', r2_score(np.array(y_test),prediction) )
```

```
/home/nauid/anaconda3/envs/tf/lib/python3.6/site-packages/ipykernel_launcher.py:4: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
  after removing the cwd from sys.path.
```

```
ann root mean absolute error:  0.12391751965552095
accuracy score:  0.9093931764105156
```

DecisionTree

Decision tree builds regression or classification models in the form of a tree structure. It breaks down a dataset into smaller and smaller subsets while at the same time an associated decision tree is incrementally developed. The final result is a tree with decision nodes and leaf nodes.

```
In [35]: from sklearn.tree import DecisionTreeRegressor
my_model = DecisionTreeRegressor()

my_model.fit(X_train, y_train)
prediction = my_model.predict(X_test)
prediction_dict['DecisionTree'] = prediction
if submit:
    submit_prediction = my_model.predict(test_processed)
    submit_prediction_dict['DecisionTree'] = submit_prediction

print('root mean absolute error: ',rmse(y_test, prediction))
print('accuracy score: ', r2_score(np.array(y_test),prediction) )

ann root mean absolute error:  0.19407309037420048
accuracy score:  0.7777580061513429
```

Xgboost

XGBoost stands for eXtreme Gradient Boosting. It is an optimized distributed gradient boosting library designed to be highly efficient, flexible and portable. It implements machine learning algorithms under the Gradient Boosting framework. XGBoost provides a parallel tree boosting (also known as GBDT, GBM) that solve many data science problems in a fast and accurate way.

```
In [36]: from xgboost import XGBRegressor
my_model = XGBRegressor(n_estimators=500, learning_rate=0.05)

my_model.fit(X_train, y_train)
prediction = my_model.predict(X_test)
prediction_dict['Xgboost'] = prediction

if submit:
    submit_prediction = my_model.predict(test_processed)
    submit_prediction_dict['Xgboost'] = submit_prediction

print('root mean absolute error: ',rmse(y_test, prediction))
print('accuracy score: ', r2_score(np.array(y_test),prediction) )

ann root mean absolute error:  0.10926796451065869
accuracy score:  0.9295499691663187
```

Lasso

Lasso (least absolute shrinkage and selection operator; also Lasso or LASSO) is a regression analysis method that performs both variable selection and regularization in order to enhance the prediction accuracy and interpretability of the statistical model it produces. Lasso was originally formulated for least squares models and this simple case reveals a substantial amount about the behavior of the estimator, including its relationship to ridge regression and best subset selection and the connections between lasso coefficient estimates and so-called soft thresholding. It also reveals that (like standard linear regression) the coefficient estimates need not be unique if covariates are collinear.

```
In [50]: from sklearn.linear_model import Lasso
my_model = Lasso(alpha=5e-3, max_iter=50000)

my_model.fit(X_train, y_train)
prediction = my_model.predict(X_test)
prediction_dict['Lasso'] = prediction

if submit:
    submit_prediction = my_model.predict(test_processed)
    submit_prediction_dict['Lasso'] = submit_prediction

print(' root mean absolute error: ', rmse(y_test, prediction))
print('accuracy score: ', r2_score(np.array(y_test), prediction) )

root mean absolute error:  0.10794617678311977
accuracy score:  0.9312440935471524
```

In the above model alpha is Constant that multiplies the L1 term. For numerical reason we cant set alpha to 0 but keeping alpha low provides good accuracy for our dataset. I have found 5e-4 provides good accuracy.

for 5e-5: root mean absolute error: 0.10973737757187135 accuracy score: 0.9289433650407954

for 1e-5: root mean absolute error: 0.11426822609093419 accuracy score: 0.9229546464396043

for 1e-3: root mean absolute error: 0.10466883446067998 accuracy score: 0.9353556969018821

for 1e-4: root mean absolute error: 0.10658498063306822 accuracy score: 0.9329671780226085

for 5e-3: root mean absolute error: 0.10794617678311977 accuracy score: 0.9312440935471524

ANN

Theory and Basics:

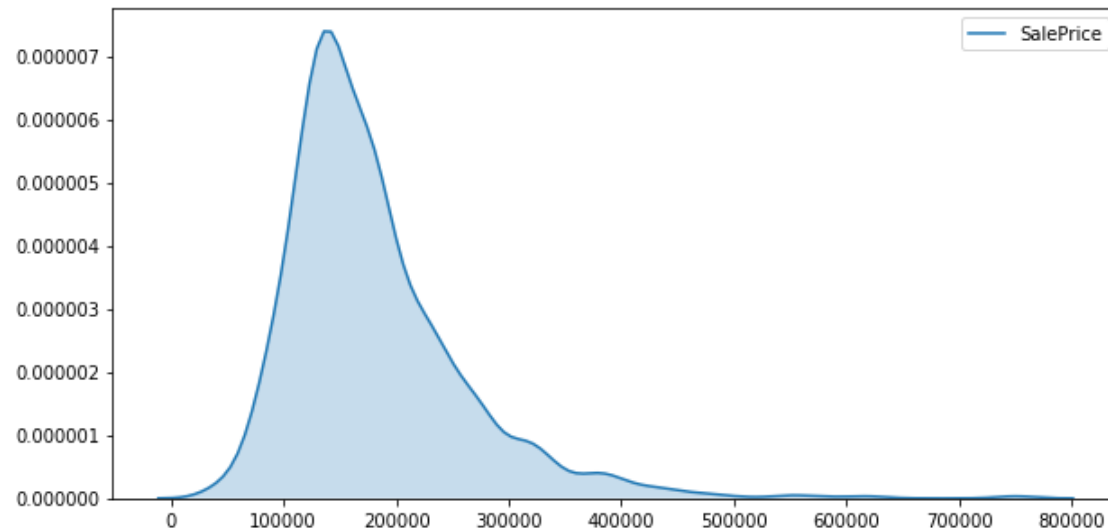
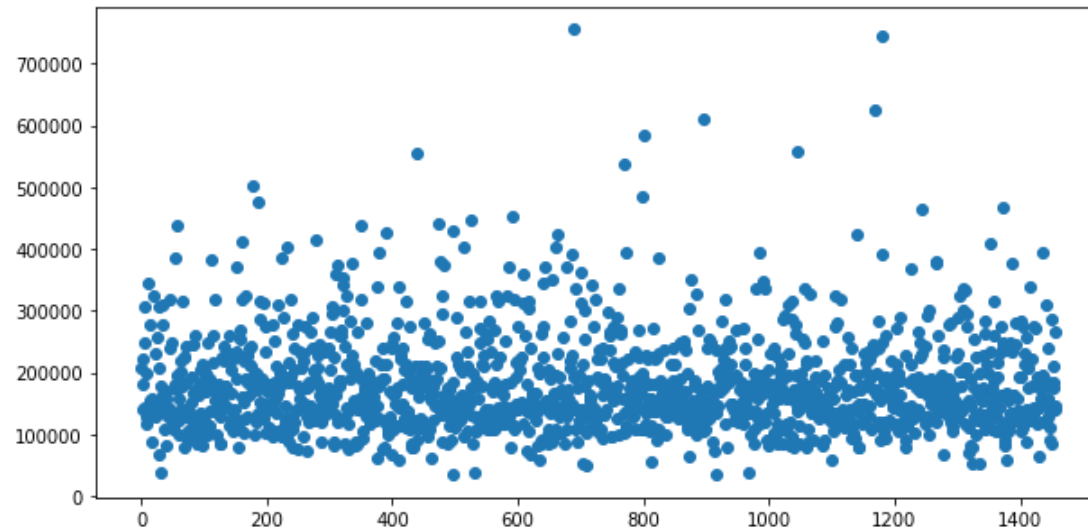
An Artificial Neural Network (ANN) is a computational model. It is based on the structure and functions of biological neural networks. It works like the way human brain processes information. ANN includes a large number of connected processing units that work together to process information. They also generate meaningful results from it.

An artificial neuron is a mathematical function conceived as a model of biological neurons, a neural network. Usually each input is separately weighted, and the sum is passed through a non-linear function known as an activation function or transfer function.

The artificial Neural network is typically organized in layers. Layers are being made up of many interconnected 'nodes' which contain an 'activation function'. A neural network may contain the following 3 layers:

- **Input layer** The purpose of the input layer is to receive as input the values of the explanatory attributes for each observation. Usually, the number of input nodes in an input layer is equal to the number of explanatory variables. 'input layer' presents the patterns to the network, which communicates to one or more 'hidden layers'. The nodes of the input layer are passive, meaning they do not change the data. They receive a single value on their input and duplicate the value to their many outputs. From the input layer, it duplicates each value and sent to all the hidden nodes.
- **Hidden Layer** The Hidden layers apply given transformations to the input values inside the network. In this, incoming arcs that go from other hidden nodes or from input nodes connected to each node. It connects with outgoing arcs to output nodes or to other hidden nodes. In hidden layer, the actual processing is done via a system of weighted 'connections'. There may be one or more hidden layers. The values entering a hidden node multiplied by weights, a set of predetermined numbers stored in the program. The weighted inputs are then added to produce a single number.
- **Output layer** The hidden layers then link to an 'output layer'. Output layer receives connections from hidden layers or from input layer. It returns an output value that corresponds to the prediction of the response variable. In classification problems, there is usually only one output node. The active nodes of the output layer combine and change the data to produce the output values.

```
In [35]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=[10,5])
plt.scatter(range(len(train)),list(train.SalePrice.values))
plt.show()
plt.figure(figsize=[10,5])
sns.kdeplot(train.SalePrice, shade= True)
plt.show()
```



In the above graph we can see that the price range is in a normal distribution. If we provide `tf.random.normal` while initializing the weight it should be more helpful for training. And this initialization should provide better validation with low amount of epoches. In my kaggle score rmse 0.123 is found through random normal while uniform distribution provided rmse 0.127 score. Again Uniform distribution takes 3 times more epoches to reach rmse score 0.127. But for uniform distribution no improvement cant found after 16000 epoch and for normal distribution no improvement can't found after 6000 epoch.

Target

By observing the span of the data and the data distribution we can conclude that logistic regression should perform well for this kind of problem. So we can safely say that starting with single neuron in a single hidden layer should perform well and we should look for simpler solution. Again from theoretical perspective single neurone and single layer ANN is nothing but a logistic regression and after adding layers and neurons we can regularize them so that they behave more like a logistic regression model and then we can tune parameter such a way so that it can handle little bit more complexity than a logistic regression. Finally my target is to make sure that it performs well as a logistic regression model and then improve it with more neuron/layers and proper tuning of parameters.

```
In [36]: # log_df = pd.DataFrame(columns=['learning_rate', 'num_steps', 'beta1','beta2','beta3', 'hidden_1' , 'hidden_2', 'hidden_3','input_dim' , 'test_rmse_score', 'test_r2_score'])
# log_df.to_csv("diffrent_training_results.csv", index=False)
```

Ann parameters

The following ANN is build with 3 hidden layers. Output dimation is 1 because its a regration problem.

```

In [37]: tf.reset_default_graph()
learning_rate = 0.1
num_steps = 6000
#for regularize weight matrix
beta1 = 0.1
beta2 = 0.00
beta3 = 0.00
beta4 = None

hidden_1 = 16
hidden_2 = 8
hidden_3 = 4
hidden_4 = None

# minimum_validation_cost is to control model saving locally
minimum_validation_cost = 0.0190000

input_dim = X_train.shape[1] # Number of features
output_dim = 1               # Because it is a regression problem
#tf graph input
X_tf = tf.placeholder("float" )
y_tf = tf.placeholder("float" )

```

Initialization of weight and bias with random values

```

In [38]: weights = {
    'w1': tf.Variable(tf.random_normal([input_dim, hidden_1])),
    'w2': tf.Variable(tf.random_normal([hidden_1, hidden_2])),
    'w3': tf.Variable(tf.random_normal([hidden_2, hidden_3])),
    'out': tf.Variable(tf.random_normal([hidden_3, output_dim]))
}
biases = {
    'b1': tf.Variable(tf.random_normal([hidden_1])),
    'b2': tf.Variable(tf.random_normal([hidden_2])),
    'b3': tf.Variable(tf.random_normal([hidden_3])),
    'out': tf.Variable(tf.random_normal([output_dim]))
}

```

WARNING:tensorflow:From /home/nauid/anaconda3/envs/tf/lib/python3.6/site-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.
Instructions for updating:
Colocations handled automatically by placer.

Model

In the dataset Sales price are non negative number so our model is expected to return positive values so as a activation function I have used relu as it gives positive values. Again relu is easy to optimize because they are similar to linear units. The only difference is that a rectified linear unit outputs zero across half its domain. Thus derivatives through a rectified linear unit remain large whenever the unit is activate. The gradients are not only large but also consistent.

```
In [39]: def ann_model(X_val):  
    # Hidden layers  
    layer_1 = tf.add(tf.matmul(X_val, weights['w1']), biases['b1'])  
    layer_1 = tf.nn.relu(layer_1)  
  
    layer_2 = tf.add(tf.matmul(layer_1, weights['w2']), biases['b2'])  
    layer_2 = tf.nn.relu(layer_2)  
  
    layer_3 = tf.add(tf.matmul(layer_2, weights['w3']), biases['b3'])  
    layer_3 = tf.nn.relu(layer_3)  
  
    # Output layer  
    layer_out = tf.add(tf.matmul(layer_3, weights['out']), biases['out'])  
  
    return layer_out
```

For optimization I have used Adam optimizer. Adam derives from phrase “adaptive moments”. Its a variant of RMSProp. I have used adam instead of RMSProp for couple of reasons. First, in Adam, momentum is incorporated directly as an estimate of the first-order moment (with exponential weighting) of the gradient. The most straightforward way to add momentum to RMSProp is to apply momentum to the rescaled gradients. The use of momentum in combination with rescaling does not have a clear theoretical motivation. Second, Adam includes bias corrections to the estimates of both the first-order moments (the momentumterm) and the (uncentered) second-order moments to account for their initializationat the origin. RMSProp also incorporates an estimate of the (uncentered) second-order moment; however, it lacks the correction factor. Thus,unlike in Adam, the RMSProp second-order moment estimate may have high bias early in training. Adam is generally regarded as being fairly robust to the choice of hyperparameters, though the learning rate sometimes needs to be changed from the suggested default. Usually default rate is .001 but for our case I have used 0.1 as it gives better optimization results.

```
In [40]: # Model Construct
model = ann_model(X_tf)

# Mean Squared Error cost function
cost = tf.reduce_mean(tf.square(y_tf - model))

# cost = tf.square(y_tf - model)
regularizer_1 = tf.nn.l2_loss(weights['w1'])
regularizer_2 = tf.nn.l2_loss(weights['w2'])
regularizer_3 = tf.nn.l2_loss(weights['w3'])
# cost = tf.reduce_mean(cost + beta1*regularizer_1 + beta2*regularizer_2 + beta3*regularizer_3)
cost = cost + beta1*regularizer_1 + beta2*regularizer_2 + beta3*regularizer_3

# Adam optimizer will update weights and biases after each step
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)

# Initialize variables
init = tf.global_variables_initializer()

# Add ops to save and restore all the variables.
saver = tf.train.Saver()
```

WARNING:tensorflow:From /home/nauid/anaconda3/envs/tf/lib/python3.6/site-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.cast instead.

Training

```
In [ ]: train_LC = []
val_LC = []
# session_var = None
```

Above train LC and val Lc variable keeps track of the learning rate so that learning curve can be drwan.

```

In [42]: def training_block(X_train,y_train, X_test,y_test):
    #reseting variables
    session_var = None
    save_path = None

    with tf.Session() as sess:
        #running initializer
        sess.run(init)

    #         minimum_validation_cost = 0.0190000
    #         tf.Session.reset(sess)
    global minimum_validation_cost
    for i in range(num_steps):
        #         sess.run(optimizer, feed_dict={X_tf:X_train, y_tf:y_train})

        if submit :
            X_train = shuffle(train_processed , random_state = i)
            y_train = shuffle(target , random_state = i)
        else:
            X_train = shuffle(X_train , random_state = i)
            y_train = shuffle(y_train , random_state = i)

        trn_cost,_ = sess.run([cost,optimizer], feed_dict={X_tf:X_train, y_tf:y_train})
        tst_cost = sess.run(cost, feed_dict={X_tf:X_test, y_tf:y_test})
        if submit :
            new_minimum_validation_cost = np.min(trn_cost)
        else:
            new_minimum_validation_cost = np.min(tst_cost)

        if (i+1)%50 == 0:
            train_LC.append(trn_cost)
            val_LC.append(tst_cost)

        if (i+1)%500 == 0:
            print("epoch no : ",i+1, "   training cost: ",trn_cost, "   validation cost: ", tst_cost, "   minimum_validation_c
ost" , minimum_validation_cost)

            if new_minimum_validation_cost < minimum_validation_cost :
                minimum_validation_cost = new_minimum_validation_cost
            #         global session_var
            #         session_var = sess
            #         Save the variables to disk.
            save_path = saver.save(sess, "model/model.ckpt")

    if bool(save_path):
        sess.close()
        print("Model saved in path: %s" % save_path)

training_block(X_train,y_train, X_test,y_test)

```

```
epoch no : 500   training cost: 4.48253   validation cost: 4.0299172   minimum_validation_cost 0.019
epoch no : 1000  training cost: 0.36521477  validation cost: 0.3668505   minimum_validation_cost 0.019
epoch no : 1500  training cost: 0.077625655  validation cost: 0.062232696  minimum_validation_cost 0.019
epoch no : 2000  training cost: 0.012304425  validation cost: 0.01430713   minimum_validation_cost 0.014317605
epoch no : 2500  training cost: 0.011156908  validation cost: 0.013151259  minimum_validation_cost 0.012850597
epoch no : 3000  training cost: 0.02404294   validation cost: 0.023661317  minimum_validation_cost 0.012850597
epoch no : 3500  training cost: 0.026674882  validation cost: 0.023832187  minimum_validation_cost 0.012850597
epoch no : 4000  training cost: 0.016393239  validation cost: 0.020106692  minimum_validation_cost 0.012850597
epoch no : 4500  training cost: 0.021710915  validation cost: 0.026027001  minimum_validation_cost 0.012850597
epoch no : 5000  training cost: 0.014663154  validation cost: 0.017669013  minimum_validation_cost 0.012850597
epoch no : 5500  training cost: 0.014944806  validation cost: 0.016238159  minimum_validation_cost 0.012850597
epoch no : 6000  training cost: 0.025703683  validation cost: 0.020001724  minimum_validation_cost 0.012850597
Model saved in path: model/model.ckpt
```

Grid search on epoch:

In the above block I have saved the model for the best validation score. As I mentioned earlier the epoch to reach the best validation accuracy is not fixed. Rather we can find it in 3 different range of epoch. The reason behind this is mostly because of random initializing of the weight and if we have fixed the seed value then it might change into only one single epoch range. But doing so we loose chance to improve our model further. Again if we want to ensemble different ANN model it woun't help when we use same seed and state. I have tried 1000+ parameters and combination from the start and used graph to visualize how to improve that but with grid search I might not get the exact idea why certain things provide good results or not and looking into every search result and graph is also too much so applying on the epoch seems to me more reasonable solution because the epoch for best validation result will be different in every run.

Trick

I have shuffled the data in every epoch and this trick improved the validation accuracy. On the other hand I didn't use batch because according to my previous experience this kind of logistic regression problem works better when its given as a whole set rather than batch or mini-batch. But if its overfitting then passing the data in a batch / mini-batch would perform better as it helps to generalize more. We can say its more like a dropout effect. And I have tried to do dropout to reduce distance of training and validation accuracy but that didn't worked well.


```

In [43]: def Prediction_block(X_test):
    with tf.Session() as sess:
        try:
            # Restore variables from disk.
            saver.restore(sess, "model/model.ckpt")
            print("Model restored.")
        except:
            print("----- available checkpoint is for different model -----")
            return
        # Check the values of the variables
        pred = sess.run(model, feed_dict={X_tf: X_test})
        prediction = pred.squeeze()
        sess.close()
        return prediction
    # print(np.exp(prediction))

prediction = Prediction_block(X_test)

pred_str = 'ANN_base_lr'+str(learning_rate)+'_beta'+str(beta1)+'-' +str(beta2)+'-' +str(beta3)+'-' +str(beta4)+'_hidden'+str(hidden_1)+'-' +str(hidden_2)+'-' +str(hidden_3)+'-' +str(hidden_4)
prediction_dict[pred_str] = prediction

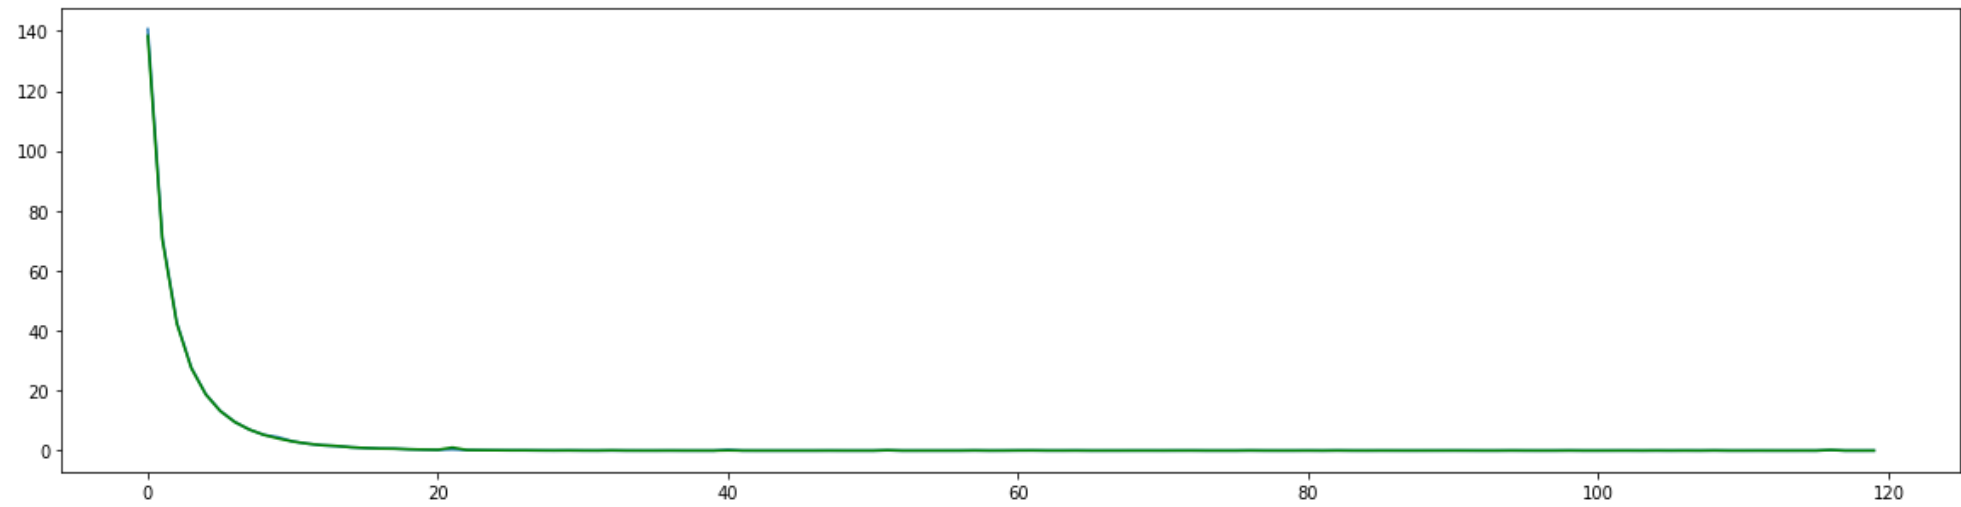
if submit:
    submit_prediction = Prediction_block(test_processed)
    submit_prediction_dict[pred_str] = submit_prediction

```

WARNING:tensorflow:From /home/nauid/anaconda3/envs/tf/lib/python3.6/site-packages/tensorflow/python/training/saver.py:1266: checkpoint_exists (from tensorflow.python.training.checkpoint_management) is deprecated and will be removed in a future version.
Instructions for updating:
Use standard file APIs to check for files with this prefix.
INFO:tensorflow:Restoring parameters from model/model.ckpt
Model restored.

Learning curve

```
In [44]: plt.figure(figsize=[20,5])  
plt.plot(train_LC)  
plt.plot(val_LC, 'g-')  
plt.show()
```



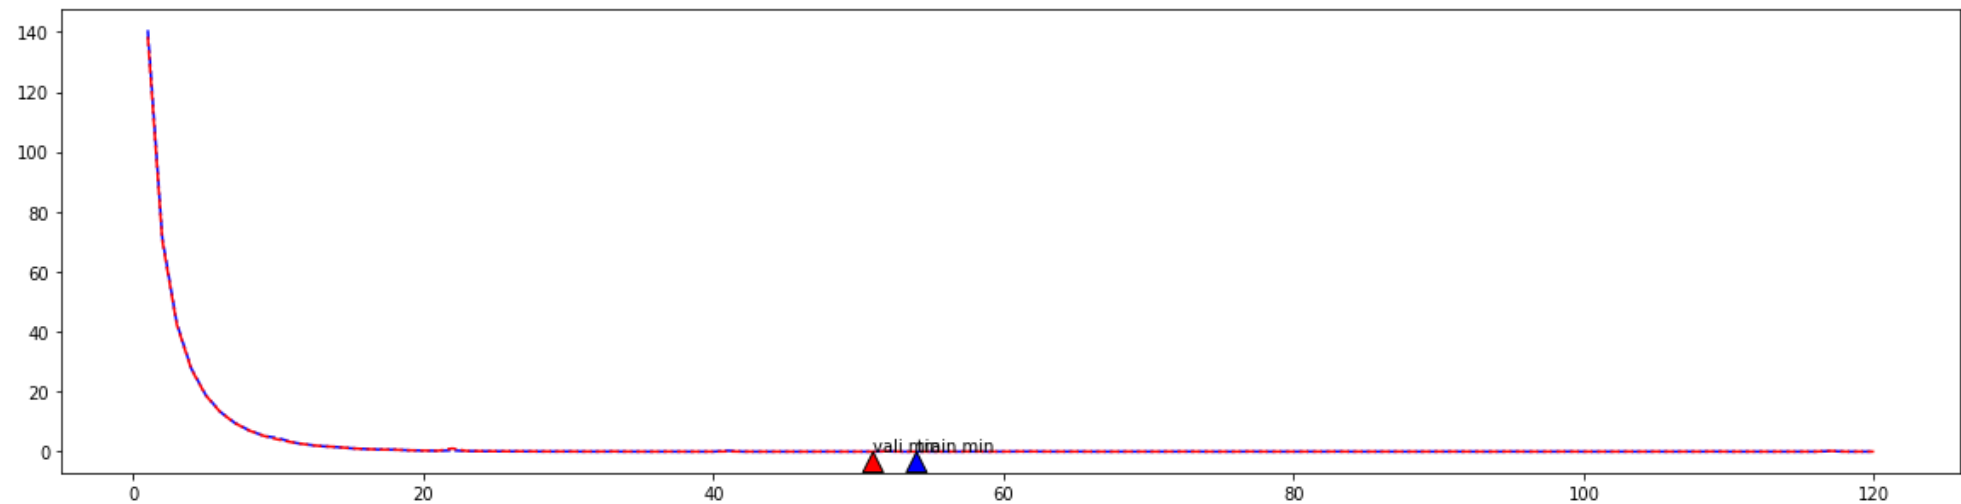
```
In [45]: def learning_curve():
xdata = list(range(1, len(train_LC)+1))
minimum = min(train_LC)

plt.figure(figsize=[20,5])
plt.plot(xdata, train_LC, 'b--')
plt.annotate('train min', xy=(xdata[train_LC.index(minimum)], minimum),
            arrowprops=dict(facecolor='blue', shrink=0.05))

minimum = min(val_LC)
plt.plot(xdata, val_LC, 'r--')
plt.annotate('vali min', xy=(xdata[val_LC.index(minimum)], minimum),
            arrowprops=dict(facecolor='red', shrink=0.05))

plt.show()

learning_curve()
```



Acuracy Score

```
In [46]: def accuracy(y_test,prediction):
    test_rmse_score = rmse(y_test, prediction)
    test_r2_score = r2_score(np.array(y_test),prediction)
    return test_rmse_score, test_r2_score

test_rmse_score, test_r2_score = accuracy(y_test,prediction)

print('ann root mean absolute error: ', test_rmse_score)
print('accuracy score: ', test_r2_score )
```

```
ann root mean absolute error:  0.10548200355447006
accuracy score:  0.9343473557962394
```

kaggle rmse:

In kaggle ranking the above ANN model provides the best rmse score and the score is 0.12324

Description on Learning curve and Accuracy:

We can observe where overfitting occurs. Overfitting actually occurs if the training loss goes under the validation loss even though the validation is still dropping. It is the sign that network is learning the patterns in the train set that are not applicable in the validation done. In a short note we can say::

Overfitting : training loss << validation loss

Underfitting : training loss >> validation loss

Just right : training loss ~ validation loss

According to this theory our both learning curve is exactly top of one another so in our case validation cost and training cost is almost same so we can say that our model is doing just the right thing. Again In validation score .1054 is impressive compared to other models.

Save score

```
In [130]: if save_score:
    log_df = pd.read_csv("different_training_results.csv")
    log_df = log_df.append({'learning_rate' : learning_rate, 'num_steps' : num_steps, 'beta1' : beta1, 'beta2' : beta2, 'beta3' :
beta3, 'beta4' : beta4, 'hidden_1' : hidden_1 , 'hidden_2' : hidden_2, 'hidden_3' : hidden_3, 'hidden_4' : hidden_4, 'input_dim'
: input_dim , 'test_rmse_score' : test_rmse_score , 'test_r2_score' : test_r2_score}, ignore_index=True)
    log_df.to_csv("different_training_results.csv", encoding='utf-8',index=False)
```

Parameters

Following segment is actually initializing different parameters. From the dataset we can see that the estimation of sale price is a regression problem and neural network used here was overfitting most of the time due to higher variance. So for making it simpler I have penalized weight matrix of hidden layers with l2 regularization. Again I have found that single hidden layer with single neuron performs well and that means the prediction model don't need to be too complex. Thus I became ensured that regularization is going to improve performance.

Cross validation

When we perform a random train-test split of our data, we assume that our examples are independent. That means that by knowing/seeing some instance will not help us understand other instances. However, that's not always the case. So to make sure if the Data is actually independent, to get more metrics and to use fine tuning my parameters on whole dataset I am performing cross validation.

```
In [49]: from sklearn.model_selection import KFold
from sklearn.model_selection import RepeatedKFold
kf = KFold(n_splits=10, shuffle=True)

kf_rmse_list = []
kf_r2_list = []

# train_processed['SalePrice'] = target.values
for train_index, test_index in kf.split(train_processed):
    X_train, X_test = train_processed.iloc[train_index] , train_processed.iloc[test_index]
    y_train, y_test = target.iloc[train_index], target.iloc[test_index]

    training_block(X_train,y_train, X_test,y_test)
    prediction = Prediction_block(X_test)
    test_rmse_score, test_r2_score = accuracy(y_test, prediction)

    kf_rmse_list.append(test_rmse_score)
    kf_r2_list.append(test_r2_score)

    print("r2 list print", kf_r2_list)
    print('rmse list print',kf_rmse_list)

print("r2 mean print", np.mean(kf_r2_list))
print('rmse mean print', np.mean(kf_rmse_list))
```

epoch no :	500	training cost:	20.568653	validation cost:	20.612467	minimum_validation_cost	0.0139135355
epoch no :	1000	training cost:	4.473949	validation cost:	4.4875116	minimum_validation_cost	0.0139135355
epoch no :	1500	training cost:	1.1820256	validation cost:	1.1890986	minimum_validation_cost	0.0139135355
epoch no :	2000	training cost:	0.30349365	validation cost:	0.3106032	minimum_validation_cost	0.0139135355
epoch no :	2500	training cost:	0.07263037	validation cost:	0.080480844	minimum_validation_cost	0.0139135355
epoch no :	3000	training cost:	0.03073879	validation cost:	0.04442262	minimum_validation_cost	0.0139135355
epoch no :	3500	training cost:	0.014154274	validation cost:	0.02467569	minimum_validation_cost	0.0139135355
epoch no :	4000	training cost:	0.07976215	validation cost:	0.05189753	minimum_validation_cost	0.0139135355
epoch no :	4500	training cost:	0.039251808	validation cost:	0.04755907	minimum_validation_cost	0.0139135355
epoch no :	5000	training cost:	0.016424252	validation cost:	0.027453901	minimum_validation_cost	0.0139135355
epoch no :	5500	training cost:	0.012508737	validation cost:	0.022914646	minimum_validation_cost	0.0139135355
epoch no :	6000	training cost:	0.014900483	validation cost:	0.020737689	minimum_validation_cost	0.0139135355

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

r2 list print [0.9553614248860762]

rmse list print [0.09271964805908306]

epoch no :	500	training cost:	11.382257	validation cost:	11.393936	minimum_validation_cost	0.0139135355
epoch no :	1000	training cost:	1.2301323	validation cost:	1.2460473	minimum_validation_cost	0.0139135355
epoch no :	1500	training cost:	0.2234553	validation cost:	0.20101692	minimum_validation_cost	0.0139135355
epoch no :	2000	training cost:	0.042567566	validation cost:	0.053197816	minimum_validation_cost	0.0139135355
epoch no :	2500	training cost:	0.022891052	validation cost:	0.03237509	minimum_validation_cost	0.0139135355
epoch no :	3000	training cost:	0.02231348	validation cost:	0.036762033	minimum_validation_cost	0.0139135355
epoch no :	3500	training cost:	0.090252094	validation cost:	0.117978	minimum_validation_cost	0.0139135355
epoch no :	4000	training cost:	0.014654718	validation cost:	0.021233875	minimum_validation_cost	0.0139135355
epoch no :	4500	training cost:	0.021118347	validation cost:	0.030047739	minimum_validation_cost	0.0139135355
epoch no :	5000	training cost:	0.020689897	validation cost:	0.023326708	minimum_validation_cost	0.0139135355
epoch no :	5500	training cost:	0.055812918	validation cost:	0.037321463	minimum_validation_cost	0.0139135355
epoch no :	6000	training cost:	0.011782758	validation cost:	0.017926298	minimum_validation_cost	0.0139135355

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

r2 list print [0.9553614248860762, 0.9508771458271212]

rmse list print [0.09271964805908306, 0.0888583273296072]

epoch no :	500	training cost:	20.266645	validation cost:	20.1501	minimum_validation_cost	0.0139135355
epoch no :	1000	training cost:	4.9247584	validation cost:	4.872145	minimum_validation_cost	0.0139135355
epoch no :	1500	training cost:	1.7483015	validation cost:	1.7037398	minimum_validation_cost	0.0139135355
epoch no :	2000	training cost:	0.75342274	validation cost:	0.71101165	minimum_validation_cost	0.0139135355
epoch no :	2500	training cost:	0.3910143	validation cost:	0.34928468	minimum_validation_cost	0.0139135355
epoch no :	3000	training cost:	0.24408573	validation cost:	0.2026229	minimum_validation_cost	0.0139135355
epoch no :	3500	training cost:	0.18683618	validation cost:	0.14548528	minimum_validation_cost	0.0139135355
epoch no :	4000	training cost:	0.16872552	validation cost:	0.12734604	minimum_validation_cost	0.0139135355
epoch no :	4500	training cost:	0.16447207	validation cost:	0.12318214	minimum_validation_cost	0.0139135355
epoch no :	5000	training cost:	0.16388221	validation cost:	0.12260775	minimum_validation_cost	0.0139135355
epoch no :	5500	training cost:	0.16384023	validation cost:	0.12256751	minimum_validation_cost	0.0139135355
epoch no :	6000	training cost:	0.16384237	validation cost:	0.122562245	minimum_validation_cost	0.0139135355

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

r2 list print [0.9553614248860762, 0.9508771458271212, 0.9543415630626542]

rmse list print [0.09271964805908306, 0.0888583273296072, 0.07449732801900573]

epoch no :	500	training cost:	0.77447015	validation cost:	0.77047086	minimum_validation_cost	0.0139135355
epoch no :	1000	training cost:	0.023457717	validation cost:	0.024699744	minimum_validation_cost	0.0139135355
epoch no :	1500	training cost:	0.011655288	validation cost:	0.012237982	minimum_validation_cost	0.012241816

epoch no :	2000	training cost:	0.011911677	validation cost:	0.012319006	minimum_validation_cost	0.011949724
epoch no :	2500	training cost:	0.012986549	validation cost:	0.013175348	minimum_validation_cost	0.011949724
epoch no :	3000	training cost:	0.013242191	validation cost:	0.013814442	minimum_validation_cost	0.011949724
epoch no :	3500	training cost:	0.013824445	validation cost:	0.013548979	minimum_validation_cost	0.011949724
epoch no :	4000	training cost:	0.014795603	validation cost:	0.01416955	minimum_validation_cost	0.011949724
epoch no :	4500	training cost:	0.03727314	validation cost:	0.078026064	minimum_validation_cost	0.011949724
epoch no :	5000	training cost:	0.13976851	validation cost:	0.07447994	minimum_validation_cost	0.011949724
epoch no :	5500	training cost:	0.015592766	validation cost:	0.014788923	minimum_validation_cost	0.011949724
epoch no :	6000	training cost:	0.023306323	validation cost:	0.0315818	minimum_validation_cost	0.011949724

Model saved in path: model/model.ckpt

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

r2 list print [0.9553614248860762, 0.9508771458271212, 0.9543415630626542, 0.9341996754951252]

rmse list print [0.09271964805908306, 0.0888583273296072, 0.07449732801900573, 0.1007255124935721]

epoch no :	500	training cost:	4.137529	validation cost:	4.1315894	minimum_validation_cost	0.011949724
epoch no :	1000	training cost:	0.90885115	validation cost:	0.9184694	minimum_validation_cost	0.011949724
epoch no :	1500	training cost:	0.32985348	validation cost:	0.34109288	minimum_validation_cost	0.011949724
epoch no :	2000	training cost:	0.18882309	validation cost:	0.20048162	minimum_validation_cost	0.011949724
epoch no :	2500	training cost:	0.16195568	validation cost:	0.1737091	minimum_validation_cost	0.011949724
epoch no :	3000	training cost:	0.15870829	validation cost:	0.17048304	minimum_validation_cost	0.011949724
epoch no :	3500	training cost:	0.15851201	validation cost:	0.17027973	minimum_validation_cost	0.011949724
epoch no :	4000	training cost:	0.15850726	validation cost:	0.1702928	minimum_validation_cost	0.011949724
epoch no :	4500	training cost:	0.1585275	validation cost:	0.17028813	minimum_validation_cost	0.011949724
epoch no :	5000	training cost:	0.1585341	validation cost:	0.17028494	minimum_validation_cost	0.011949724
epoch no :	5500	training cost:	0.15850429	validation cost:	0.1702753	minimum_validation_cost	0.011949724
epoch no :	6000	training cost:	0.15850051	validation cost:	0.1702815	minimum_validation_cost	0.011949724

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

r2 list print [0.9553614248860762, 0.9508771458271212, 0.9543415630626542, 0.9341996754951252, 0.9504439704044774]

rmse list print [0.09271964805908306, 0.0888583273296072, 0.07449732801900573, 0.1007255124935721, 0.09159195884801004]

epoch no :	500	training cost:	0.18420246	validation cost:	0.17968395	minimum_validation_cost	0.011949724
epoch no :	1000	training cost:	0.16007997	validation cost:	0.15593112	minimum_validation_cost	0.011949724
epoch no :	1500	training cost:	0.160098	validation cost:	0.15596539	minimum_validation_cost	0.011949724
epoch no :	2000	training cost:	0.16007857	validation cost:	0.15593415	minimum_validation_cost	0.011949724
epoch no :	2500	training cost:	0.16015163	validation cost:	0.15594558	minimum_validation_cost	0.011949724
epoch no :	3000	training cost:	0.16008155	validation cost:	0.15593125	minimum_validation_cost	0.011949724
epoch no :	3500	training cost:	0.16009273	validation cost:	0.15595405	minimum_validation_cost	0.011949724
epoch no :	4000	training cost:	0.16012664	validation cost:	0.15593708	minimum_validation_cost	0.011949724
epoch no :	4500	training cost:	0.16007891	validation cost:	0.1559313	minimum_validation_cost	0.011949724
epoch no :	5000	training cost:	0.1601019	validation cost:	0.15595824	minimum_validation_cost	0.011949724
epoch no :	5500	training cost:	0.16012958	validation cost:	0.15597698	minimum_validation_cost	0.011949724
epoch no :	6000	training cost:	0.1601085	validation cost:	0.15594654	minimum_validation_cost	0.011949724

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

r2 list print [0.9553614248860762, 0.9508771458271212, 0.9543415630626542, 0.9341996754951252, 0.9504439704044774, 0.929886933228908]

rmse list print [0.09271964805908306, 0.0888583273296072, 0.07449732801900573, 0.1007255124935721, 0.09159195884801004, 0.10455838190668743]

epoch no :	500	training cost:	0.17587255	validation cost:	0.19876517	minimum_validation_cost	0.011949724
epoch no :	1000	training cost:	0.15733822	validation cost:	0.18056849	minimum_validation_cost	0.011949724
epoch no :	1500	training cost:	0.157345	validation cost:	0.18057846	minimum_validation_cost	0.011949724

epoch no :	2000	training cost:	0.15734199	validation cost:	0.18057579	minimum_validation_cost	0.011949724
epoch no :	2500	training cost:	0.15734023	validation cost:	0.18057136	minimum_validation_cost	0.011949724
epoch no :	3000	training cost:	0.15734643	validation cost:	0.18057187	minimum_validation_cost	0.011949724
epoch no :	3500	training cost:	0.15733868	validation cost:	0.18056941	minimum_validation_cost	0.011949724
epoch no :	4000	training cost:	0.15737869	validation cost:	0.18057455	minimum_validation_cost	0.011949724
epoch no :	4500	training cost:	0.15742067	validation cost:	0.18062714	minimum_validation_cost	0.011949724
epoch no :	5000	training cost:	0.1573638	validation cost:	0.18058623	minimum_validation_cost	0.011949724
epoch no :	5500	training cost:	0.15738843	validation cost:	0.18061881	minimum_validation_cost	0.011949724
epoch no :	6000	training cost:	0.15735286	validation cost:	0.18060046	minimum_validation_cost	0.011949724

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

r2 list print [0.9553614248860762, 0.9508771458271212, 0.9543415630626542, 0.9341996754951252, 0.9504439704044774, 0.929886933228908, 0.9382959161998816]

rmse list print [0.09271964805908306, 0.0888583273296072, 0.07449732801900573, 0.1007255124935721, 0.09159195884801004, 0.10455838190668743, 0.10551438239716186]

epoch no :	500	training cost:	0.19058843	validation cost:	0.19598773	minimum_validation_cost	0.011949724
epoch no :	1000	training cost:	0.12674819	validation cost:	0.090106644	minimum_validation_cost	0.011949724
epoch no :	1500	training cost:	0.010776039	validation cost:	0.015417905	minimum_validation_cost	0.011949724
epoch no :	2000	training cost:	0.011808951	validation cost:	0.016529392	minimum_validation_cost	0.011949724
epoch no :	2500	training cost:	0.012140089	validation cost:	0.017381266	minimum_validation_cost	0.011949724
epoch no :	3000	training cost:	0.020829326	validation cost:	0.01790555	minimum_validation_cost	0.011949724
epoch no :	3500	training cost:	0.01308694	validation cost:	0.016618628	minimum_validation_cost	0.011949724
epoch no :	4000	training cost:	0.017044349	validation cost:	0.021237336	minimum_validation_cost	0.011949724
epoch no :	4500	training cost:	0.013240088	validation cost:	0.016861184	minimum_validation_cost	0.011949724
epoch no :	5000	training cost:	0.013989	validation cost:	0.017500032	minimum_validation_cost	0.011949724
epoch no :	5500	training cost:	0.014921311	validation cost:	0.018324185	minimum_validation_cost	0.011949724
epoch no :	6000	training cost:	0.01547069	validation cost:	0.018815495	minimum_validation_cost	0.011949724

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

r2 list print [0.9553614248860762, 0.9508771458271212, 0.9543415630626542, 0.9341996754951252, 0.9504439704044774, 0.929886933228908, 0.9382959161998816, 0.932418870844409]

rmse list print [0.09271964805908306, 0.0888583273296072, 0.07449732801900573, 0.1007255124935721, 0.09159195884801004, 0.10455838190668743, 0.10551438239716186, 0.10189308411814846]

epoch no :	500	training cost:	14.345135	validation cost:	14.309958	minimum_validation_cost	0.011949724
epoch no :	1000	training cost:	1.9079558	validation cost:	1.9146067	minimum_validation_cost	0.011949724
epoch no :	1500	training cost:	0.45495993	validation cost:	0.37160826	minimum_validation_cost	0.011949724
epoch no :	2000	training cost:	0.10315405	validation cost:	0.10432874	minimum_validation_cost	0.011949724
epoch no :	2500	training cost:	0.048060406	validation cost:	0.07186215	minimum_validation_cost	0.011949724
epoch no :	3000	training cost:	0.022855159	validation cost:	0.021485604	minimum_validation_cost	0.011949724
epoch no :	3500	training cost:	0.012944125	validation cost:	0.012365152	minimum_validation_cost	0.010917494
epoch no :	4000	training cost:	0.011641011	validation cost:	0.010896233	minimum_validation_cost	0.010901678
epoch no :	4500	training cost:	0.011573468	validation cost:	0.0106407385	minimum_validation_cost	0.010396154
epoch no :	5000	training cost:	0.012472791	validation cost:	0.011355851	minimum_validation_cost	0.010396154
epoch no :	5500	training cost:	0.012279389	validation cost:	0.011330111	minimum_validation_cost	0.010396154
epoch no :	6000	training cost:	0.0148251755	validation cost:	0.01329639	minimum_validation_cost	0.010396154

Model saved in path: model/model.ckpt

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

r2 list print [0.9553614248860762, 0.9508771458271212, 0.9543415630626542, 0.9341996754951252, 0.9504439704044774, 0.929886933228908, 0.9382959161998816, 0.932418870844409, 0.9366822053541947]

rmse list print [0.09271964805908306, 0.0888583273296072, 0.07449732801900573, 0.1007255124935721, 0.09159195884801004, 0.10455838190668743, 0.10551438239716186, 0.10189308411814846]

```

8190668743, 0.10551438239716186, 0.10189308411814846, 0.09402134306063017]
epoch no : 500    training cost: 2.672395    validation cost: 2.670042    minimum_validation_cost 0.010396154
epoch no : 1000   training cost: 0.13111192   validation cost: 0.13321456   minimum_validation_cost 0.010396154
epoch no : 1500   training cost: 0.011619842   validation cost: 0.019234128   minimum_validation_cost 0.010396154
epoch no : 2000   training cost: 0.009893612   validation cost: 0.017380662   minimum_validation_cost 0.010396154
epoch no : 2500   training cost: 0.010010591   validation cost: 0.017384036   minimum_validation_cost 0.010396154
epoch no : 3000   training cost: 0.015444754   validation cost: 0.02303415    minimum_validation_cost 0.010396154
epoch no : 3500   training cost: 0.015875159   validation cost: 0.028511856   minimum_validation_cost 0.010396154
epoch no : 4000   training cost: 0.012220575   validation cost: 0.018918835   minimum_validation_cost 0.010396154
epoch no : 4500   training cost: 0.013949025   validation cost: 0.022416791   minimum_validation_cost 0.010396154
epoch no : 5000   training cost: 0.013269598   validation cost: 0.020984132   minimum_validation_cost 0.010396154
epoch no : 5500   training cost: 0.013208384   validation cost: 0.019760473   minimum_validation_cost 0.010396154
epoch no : 6000   training cost: 0.026704477   validation cost: 0.031736426   minimum_validation_cost 0.010396154
INFO:tensorflow:Restoring parameters from model/model.ckpt
Model restored.
r2 list print [0.9553614248860762, 0.9508771458271212, 0.9543415630626542, 0.9341996754951252, 0.9504439704044774, 0.929886933228
908, 0.9382959161998816, 0.932418870844409, 0.9366822053541947, 0.9255999030483704]
rmse list print [0.09271964805908306, 0.0888583273296072, 0.07449732801900573, 0.1007255124935721, 0.09159195884801004, 0.1045583
8190668743, 0.10551438239716186, 0.10189308411814846, 0.09402134306063017, 0.11009827566438594]
r2 mean print 0.9408107608351217
rmse mean print 0.09644782418962919

```

Observation

In the cross validation section we can see that 10 fold cross validation on our best ANN model provides similar rmse to 80-20 split rmse score. So we can rely on 80-20 split on this dataset. Thus we can say that the data in the dataset is independent.

Observing Few Other well performed ANN models

In this section We are observing the few other models and their learning curve. After that some of them will be used for Ensemble learning section for further improvement. In this model I have only changed the size of hidden layer, amount of neuron in each hidden layers , number of steps and learning rates. Rest of the part is same as the ANN described above.

ANN with 4 layers

Initialization of models

```
In [100]: tf.reset_default_graph()
def weight_bais():
    global weights, biases
    weights = None
    biases = None

    weights = {
        'w1': tf.Variable(tf.random_normal([input_dim, hidden_1])),
        'w2': tf.Variable(tf.random_normal([hidden_1, hidden_2])),
        'w3': tf.Variable(tf.random_normal([hidden_2, hidden_3])),
        'w4': tf.Variable(tf.random_normal([hidden_3, hidden_4])),
        'out': tf.Variable(tf.random_normal([hidden_4, output_dim]))
    }
    biases = {
        'b1': tf.Variable(tf.random_normal([hidden_1])),
        'b2': tf.Variable(tf.random_normal([hidden_2])),
        'b3': tf.Variable(tf.random_normal([hidden_3])),
        'b4': tf.Variable(tf.random_normal([hidden_4])),
        'out': tf.Variable(tf.random_normal([output_dim]))
    }
```

```
In [101]: def ann_model(X_val):
    # Hidden layers
    layer_1 = tf.add(tf.matmul(X_val, weights['w1']), biases['b1'])
    layer_1 = tf.nn.relu(layer_1)

    layer_2 = tf.add(tf.matmul(layer_1, weights['w2']), biases['b2'])
    layer_2 = tf.nn.relu(layer_2)

    layer_3 = tf.add(tf.matmul(layer_2, weights['w3']), biases['b3'])
    layer_3 = tf.nn.relu(layer_3)

    layer_4 = tf.add(tf.matmul(layer_3, weights['w4']), biases['b4'])
    layer_4 = tf.nn.relu(layer_4)

    # Output layer
    layer_out = tf.add(tf.matmul(layer_4, weights['out']), biases['out'])

    return layer_out
```

```

In [102]: regularizer_4 = None
def miscellaneous_initialization():
    global model, cost , regularizer_1 , regularizer_2 ,regularizer_3, regularizer_4, optimizer , init , saver
    # Model Construct
    model = ann_model(X_tf)

    # Mean Squared Error cost function
    cost = tf.reduce_mean(tf.square(y_tf - model))

    # cost = tf.square(y_tf - model)
    regularizer_1 = tf.nn.l2_loss(weights['w1'])
    regularizer_2 = tf.nn.l2_loss(weights['w2'])
    regularizer_3 = tf.nn.l2_loss(weights['w3'])
    regularizer_4 = tf.nn.l2_loss(weights['w4'])
    # cost = tf.reduce_mean(cost + beta1*regularizer_1 + beta2*regularizer_2 + beta3*regularizer_3)
    cost = cost + beta1*regularizer_1 + beta2*regularizer_2 + beta3*regularizer_3 + beta4*regularizer_4

    # Adam optimizer will update weights and biases after each step
    optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)

    # Initialize variables
    init = tf.global_variables_initializer()

    # Add ops to save and restore all the variables.
    saver = tf.train.Saver()

```

Training

ANN 1

In this section changed variables are

- learning rate = .01

layer name	Neuron	value of beta for l2 regularization
1st hidden layer	76 Neuron	.1
2nd hidden layer	48 Neuron	.05
3rd hidden layer	32 Neuron	0
4th hidden layer	16 Neuron	0

```
In [253]: tf.reset_default_graph()
learning_rate = 0.1
num_steps = 25000
#for regularize weight matrix
beta1 = 0.1
beta2 = 0.05
beta3 = 0.00
beta4 = 0.0

hidden_1 = 76
hidden_2 = 48
hidden_3 = 32
hidden_4 = 16

minimum_validation_cost = .02101000

input_dim = X_train.shape[1] # Number of features
output_dim = 1 # Because it is a regression problem

#tf graph input
X_tf = tf.placeholder("float" )
y_tf = tf.placeholder("float" )

weight_bais()
miscellaneous_initialization()
train_LC = []
val_LC = []
```

```

In [254]: training_block(X_train,y_train, X_test,y_test)
prediction = Prediction_block(X_test)
test_rmse_score, test_r2_score = accuracy(y_test,prediction)

print('ann root mean absolute error: ', test_rmse_score)
print('accuracy score: ', test_r2_score )
learning_curve()

pred_str = 'ANN_lr'+str(learning_rate)+'_beta'+str(beta1)+'-'+str(beta2)+'-'+str(beta3)+'-'+str(beta4)+'_hidden'+str(hidden_1)+'-'+str(hidden_2)+'-'+str(hidden_3)+'-'+str(hidden_4)
prediction_dict[pred_str] = prediction


if submit:
    submit_prediction = Prediction_block(test_processed)
    submit_prediction_dict[pred_str] = submit_prediction
# Data Save
if save_score:
    log_df = pd.read_csv("diffrent_training_results.csv")
    log_df = log_df.append({'learning_rate' : learning_rate, 'num_steps' : num_steps, 'beta1' : beta1, 'beta2' : beta2, 'beta3' : beta3, 'beta4' : beta4, 'hidden_1' : hidden_1 , 'hidden_2' : hidden_2, 'hidden_3' : hidden_3, 'hidden_4' : hidden_4, 'input_dim' : input_dim , 'test_rmse_score' : test_rmse_score , 'test_r2_score' : test_r2_score}, ignore_index=True)
    log_df.to_csv("diffrent_training_results.csv", encoding='utf-8',index=False)

```

epoch no :	500	training cost:	1948.0289	validation cost:	1963.5165	minimum_validation_cost	0.02101
epoch no :	1000	training cost:	1789.3033	validation cost:	1804.5315	minimum_validation_cost	0.02101
epoch no :	1500	training cost:	1631.0383	validation cost:	1646.2167	minimum_validation_cost	0.02101
epoch no :	2000	training cost:	1484.1322	validation cost:	1499.1335	minimum_validation_cost	0.02101
epoch no :	2500	training cost:	1351.5776	validation cost:	1365.8958	minimum_validation_cost	0.02101
epoch no :	3000	training cost:	1231.5188	validation cost:	1244.1323	minimum_validation_cost	0.02101
epoch no :	3500	training cost:	1121.0094	validation cost:	1129.8462	minimum_validation_cost	0.02101
epoch no :	4000	training cost:	1016.4562	validation cost:	1022.328	minimum_validation_cost	0.02101
epoch no :	4500	training cost:	915.9035	validation cost:	919.5465	minimum_validation_cost	0.02101
epoch no :	5000	training cost:	818.86554	validation cost:	821.1494	minimum_validation_cost	0.02101
epoch no :	5500	training cost:	725.55963	validation cost:	729.7903	minimum_validation_cost	0.02101
epoch no :	6000	training cost:	636.4901	validation cost:	640.9389	minimum_validation_cost	0.02101
epoch no :	6500	training cost:	552.2778	validation cost:	555.0504	minimum_validation_cost	0.02101
epoch no :	7000	training cost:	473.5905	validation cost:	474.14694	minimum_validation_cost	0.02101
epoch no :	7500	training cost:	401.07443	validation cost:	401.0503	minimum_validation_cost	0.02101
epoch no :	8000	training cost:	335.22263	validation cost:	335.2226	minimum_validation_cost	0.02101
epoch no :	8500	training cost:	276.37018	validation cost:	276.4013	minimum_validation_cost	0.02101
epoch no :	9000	training cost:	224.65007	validation cost:	224.68013	minimum_validation_cost	0.02101
epoch no :	9500	training cost:	179.96939	validation cost:	179.98999	minimum_validation_cost	0.02101
epoch no :	10000	training cost:	142.05312	validation cost:	142.03314	minimum_validation_cost	0.02101
epoch no :	10500	training cost:	110.439835	validation cost:	110.40212	minimum_validation_cost	0.02101
epoch no :	11000	training cost:	84.54955	validation cost:	84.55469	minimum_validation_cost	0.02101
epoch no :	11500	training cost:	63.74702	validation cost:	63.77253	minimum_validation_cost	0.02101
epoch no :	12000	training cost:	47.36451	validation cost:	47.38588	minimum_validation_cost	0.02101
epoch no :	12500	training cost:	34.71665	validation cost:	34.722527	minimum_validation_cost	0.02101
epoch no :	13000	training cost:	25.131432	validation cost:	25.155582	minimum_validation_cost	0.02101
epoch no :	13500	training cost:	17.955591	validation cost:	17.975864	minimum_validation_cost	0.02101
epoch no :	14000	training cost:	12.609146	validation cost:	12.616015	minimum_validation_cost	0.02101
epoch no :	14500	training cost:	8.702005	validation cost:	8.7001095	minimum_validation_cost	0.02101
epoch no :	15000	training cost:	5.90036	validation cost:	5.8999863	minimum_validation_cost	0.02101
epoch no :	15500	training cost:	3.8992198	validation cost:	3.8999383	minimum_validation_cost	0.02101
epoch no :	16000	training cost:	2.5155318	validation cost:	2.5166104	minimum_validation_cost	0.02101
epoch no :	16500	training cost:	1.5840064	validation cost:	1.585098	minimum_validation_cost	0.02101
epoch no :	17000	training cost:	0.9745886	validation cost:	0.9757318	minimum_validation_cost	0.02101
epoch no :	17500	training cost:	0.5837551	validation cost:	0.5853919	minimum_validation_cost	0.02101
epoch no :	18000	training cost:	0.34044278	validation cost:	0.34229845	minimum_validation_cost	0.02101
epoch no :	18500	training cost:	0.19107433	validation cost:	0.19231078	minimum_validation_cost	0.02101
epoch no :	19000	training cost:	0.106256224	validation cost:	0.10772059	minimum_validation_cost	0.02101
epoch no :	19500	training cost:	0.060021322	validation cost:	0.062279984	minimum_validation_cost	0.02101
epoch no :	20000	training cost:	0.034057993	validation cost:	0.03559231	minimum_validation_cost	0.02101
epoch no :	20500	training cost:	0.022776678	validation cost:	0.024662416	minimum_validation_cost	0.02101
epoch no :	21000	training cost:	0.026222728	validation cost:	0.026665546	minimum_validation_cost	0.019980568
epoch no :	21500	training cost:	0.018489484	validation cost:	0.01983606	minimum_validation_cost	0.018412706
epoch no :	22000	training cost:	0.01716589	validation cost:	0.018984884	minimum_validation_cost	0.018258983
epoch no :	22500	training cost:	0.019204231	validation cost:	0.020943124	minimum_validation_cost	0.018142112
epoch no :	23000	training cost:	0.023870615	validation cost:	0.031735398	minimum_validation_cost	0.018142112
epoch no :	23500	training cost:	0.017715417	validation cost:	0.022381995	minimum_validation_cost	0.018142112
epoch no :	24000	training cost:	0.019991433	validation cost:	0.023535162	minimum_validation_cost	0.018142112
epoch no :	24500	training cost:	0.02650885	validation cost:	0.025393497	minimum_validation_cost	0.018142112
epoch no :	25000	training cost:	0.061272897	validation cost:	0.2311845	minimum_validation_cost	0.018142112

Model saved in path: model/model.ckpt

INFO:tensorflow:Restoring parameters from model/model.ckpt
Model restored.
ann root mean absolute error: 0.11033837928515242
accuracy score: 0.9281629180762654



ANN 2

In this section changed variables are

- learning rate = .05

layer name	Neuron	value of beta for l2 regularization
1st hidden layer	8 Neuron	.005
2nd hidden layer	32 Neuron	.1
3rd hidden layer	16 Neuron	0.05
4th hidden layer	8 Neuron	0


```
In [55]: tf.reset_default_graph()
learning_rate = 0.05
num_steps = 25000
#for regularize weight matrix
beta1 = 0.005
beta2 = 0.1
beta3 = 0.05
beta4 = 0.0

hidden_1 = 8
hidden_2 = 32
hidden_3 = 16
hidden_4 = 8

minimum_validation_cost = 0.02101000
#tf graph input
X_tf = tf.placeholder("float" )
y_tf = tf.placeholder("float" )

weight_bais()
miscellaneous_initialization()
train_LC = []
val_LC = []
```

```

In [56]: training_block(X_train,y_train, X_test,y_test)
prediction = Prediction_block(X_test)

test_rmse_score, test_r2_score = accuracy(y_test,prediction)

print('ann root mean absolute error: ', test_rmse_score)
print('accuracy score: ', test_r2_score )
learning_curve()

pred_str = 'ANN_lr'+str(learning_rate)+'_beta'+str(beta1)+'-'+str(beta2)+'-'+str(beta3)+'-'+str(beta4)+'_hidden'+str(hidden_1)+'-'+str(hidden_2)+'-'+str(hidden_3)+'-'+str(hidden_4)
prediction_dict[pred_str] = prediction


if submit:
    submit_prediction = Prediction_block(test_processed)
    submit_prediction_dict[pred_str] = submit_prediction
# Data Save
if save_score:
    log_df = pd.read_csv("different_training_results.csv")
    log_df = log_df.append({'learning_rate' : learning_rate, 'num_steps' : num_steps, 'beta1' : beta1, 'beta2' : beta2, 'beta3' : beta3, 'beta4' : beta4, 'hidden_1' : hidden_1 , 'hidden_2' : hidden_2, 'hidden_3' : hidden_3, 'hidden_4' : hidden_4, 'input_dim' : input_dim , 'test_rmse_score' : test_rmse_score , 'test_r2_score' : test_r2_score}, ignore_index=True)
    log_df.to_csv("different_training_results.csv", encoding='utf-8',index=False)

```

epoch no :	500	training cost:	98.96163	validation cost:	101.16655	minimum_validation_cost	0.02101
epoch no :	1000	training cost:	45.49598	validation cost:	46.24627	minimum_validation_cost	0.02101
epoch no :	1500	training cost:	31.685268	validation cost:	31.812305	minimum_validation_cost	0.02101
epoch no :	2000	training cost:	28.494917	validation cost:	28.509579	minimum_validation_cost	0.02101
epoch no :	2500	training cost:	26.541996	validation cost:	26.54794	minimum_validation_cost	0.02101
epoch no :	3000	training cost:	24.718468	validation cost:	24.729013	minimum_validation_cost	0.02101
epoch no :	3500	training cost:	22.942047	validation cost:	22.95431	minimum_validation_cost	0.02101
epoch no :	4000	training cost:	21.18252	validation cost:	21.19542	minimum_validation_cost	0.02101
epoch no :	4500	training cost:	19.437643	validation cost:	19.448608	minimum_validation_cost	0.02101
epoch no :	5000	training cost:	17.709394	validation cost:	17.720444	minimum_validation_cost	0.02101
epoch no :	5500	training cost:	15.999163	validation cost:	16.01024	minimum_validation_cost	0.02101
epoch no :	6000	training cost:	14.31156	validation cost:	14.322832	minimum_validation_cost	0.02101
epoch no :	6500	training cost:	12.659315	validation cost:	12.677395	minimum_validation_cost	0.02101
epoch no :	7000	training cost:	11.062086	validation cost:	11.081362	minimum_validation_cost	0.02101
epoch no :	7500	training cost:	9.52519	validation cost:	9.536628	minimum_validation_cost	0.02101
epoch no :	8000	training cost:	8.041063	validation cost:	8.0620365	minimum_validation_cost	0.02101
epoch no :	8500	training cost:	6.627181	validation cost:	6.6679873	minimum_validation_cost	0.02101
epoch no :	9000	training cost:	5.356043	validation cost:	5.36625	minimum_validation_cost	0.02101
epoch no :	9500	training cost:	4.2130976	validation cost:	4.231335	minimum_validation_cost	0.02101
epoch no :	10000	training cost:	3.2202468	validation cost:	3.2339666	minimum_validation_cost	0.02101
epoch no :	10500	training cost:	2.4022565	validation cost:	2.420408	minimum_validation_cost	0.02101
epoch no :	11000	training cost:	1.7369599	validation cost:	1.7432039	minimum_validation_cost	0.02101
epoch no :	11500	training cost:	1.2229878	validation cost:	1.2248117	minimum_validation_cost	0.02101
epoch no :	12000	training cost:	0.8650217	validation cost:	0.8522291	minimum_validation_cost	0.02101
epoch no :	12500	training cost:	0.5642521	validation cost:	0.5668021	minimum_validation_cost	0.02101
epoch no :	13000	training cost:	0.3693418	validation cost:	0.37263116	minimum_validation_cost	0.02101
epoch no :	13500	training cost:	0.23565769	validation cost:	0.23976843	minimum_validation_cost	0.02101
epoch no :	14000	training cost:	0.15317412	validation cost:	0.15297893	minimum_validation_cost	0.02101
epoch no :	14500	training cost:	0.089964926	validation cost:	0.095126145	minimum_validation_cost	0.02101
epoch no :	15000	training cost:	0.05444552	validation cost:	0.060352698	minimum_validation_cost	0.02101
epoch no :	15500	training cost:	0.033405006	validation cost:	0.04072603	minimum_validation_cost	0.02101
epoch no :	16000	training cost:	0.021379678	validation cost:	0.02719137	minimum_validation_cost	0.02101
epoch no :	16500	training cost:	0.017818375	validation cost:	0.02784073	minimum_validation_cost	0.02101
epoch no :	17000	training cost:	0.01326775	validation cost:	0.021916628	minimum_validation_cost	0.020162757
epoch no :	17500	training cost:	0.017692542	validation cost:	0.026436877	minimum_validation_cost	0.018723719
epoch no :	18000	training cost:	0.012610511	validation cost:	0.019203333	minimum_validation_cost	0.018243955
epoch no :	18500	training cost:	0.011309475	validation cost:	0.018362269	minimum_validation_cost	0.018199008
epoch no :	19000	training cost:	0.017165741	validation cost:	0.026090968	minimum_validation_cost	0.018195953
epoch no :	19500	training cost:	0.017472688	validation cost:	0.025563277	minimum_validation_cost	0.018068379
epoch no :	20000	training cost:	0.01170708	validation cost:	0.018411675	minimum_validation_cost	0.017956132
epoch no :	20500	training cost:	0.017795723	validation cost:	0.023501392	minimum_validation_cost	0.017956132
epoch no :	21000	training cost:	0.013715736	validation cost:	0.02364869	minimum_validation_cost	0.017956132
epoch no :	21500	training cost:	0.01588199	validation cost:	0.024066374	minimum_validation_cost	0.01679543
epoch no :	22000	training cost:	0.013990714	validation cost:	0.019897664	minimum_validation_cost	0.01679543
epoch no :	22500	training cost:	0.026997136	validation cost:	0.021226006	minimum_validation_cost	0.01679543
epoch no :	23000	training cost:	0.014782673	validation cost:	0.019962754	minimum_validation_cost	0.01679543
epoch no :	23500	training cost:	0.02022815	validation cost:	0.026848074	minimum_validation_cost	0.01679543
epoch no :	24000	training cost:	0.016962774	validation cost:	0.021317616	minimum_validation_cost	0.01679543
epoch no :	24500	training cost:	0.024479048	validation cost:	0.02893193	minimum_validation_cost	0.01679543
epoch no :	25000	training cost:	0.021377431	validation cost:	0.033468053	minimum_validation_cost	0.01679543

Model saved in path: model/model.ckpt

INFO:tensorflow:Restoring parameters from model/model.ckpt
Model restored.
ann root mean absolute error: 0.10818300213571848
accuracy score: 0.9309420725912334



ANN 3

- learning rate = .05

layer name	Neuron	value of beta for l2 regularization
1st hidden layer	16 Neuron	.1
2nd hidden layer	8 Neuron	.0
3rd hidden layer	4 Neuron	0.0
4th hidden layer	2 Neuron	0

```
In [63]: tf.reset_default_graph()
learning_rate = 0.05
num_steps = 15000
#for regularize weight matrix
beta1 = 0.1
beta2 = 0.0
beta3 = 0.00
beta4 = 0.0

hidden_1 = 16
hidden_2 = 8
hidden_3 = 4
hidden_4 = 2

minimum_validation_cost = 0.01901000

#tf graph input
X_tf = tf.placeholder("float" )
y_tf = tf.placeholder("float" )

weight_bais()
miscellaneous_initialization()
train_LC = []
val_LC = []
```

```

In [64]: training_block(X_train,y_train, X_test,y_test)
prediction = Prediction_block(X_test)
test_rmse_score, test_r2_score = accuracy(y_test,prediction)

print('ann root mean absolute error: ', test_rmse_score)
print('accuracy score: ', test_r2_score )
learning_curve()

pred_str = 'ANN_lr'+str(learning_rate)+'_beta'+str(beta1)+'-'+str(beta2)+'-'+str(beta3)+'-'+str(beta4)+'_hidden'+str(hidden_1)+'-'+str(hidden_2)+'-'+str(hidden_3)+'-'+str(hidden_4)
prediction_dict[pred_str] = prediction

if submit:
    submit_prediction = Prediction_block(test_processed)
    submit_prediction_dict[pred_str] = submit_prediction
# Data Save
if save_score:
    log_df = pd.read_csv("diffrent_training_results.csv")
    log_df = log_df.append({'learning_rate' : learning_rate, 'num_steps' : num_steps, 'beta1' : beta1, 'beta2' : beta2, 'beta3' : beta3, 'beta4' : beta4, 'hidden_1' : hidden_1 , 'hidden_2' : hidden_2, 'hidden_3' : hidden_3, 'hidden_4' : hidden_4, 'input_dim' : input_dim , 'test_rmse_score' : test_rmse_score , 'test_r2_score' : test_r2_score}, ignore_index=True)
    log_df.to_csv("diffrent_training_results.csv", encoding='utf-8',index=False)

```

epoch no :	500	training cost:	0.06750769	validation cost:	0.07170834	minimum_validation_cost	0.01901
epoch no :	1000	training cost:	0.024438005	validation cost:	0.027701719	minimum_validation_cost	0.01901
epoch no :	1500	training cost:	0.10088402	validation cost:	0.03532496	minimum_validation_cost	0.015040747
epoch no :	2000	training cost:	0.011049605	validation cost:	0.013185406	minimum_validation_cost	0.012483893
epoch no :	2500	training cost:	0.011799688	validation cost:	0.013729342	minimum_validation_cost	0.012483893
epoch no :	3000	training cost:	0.013513778	validation cost:	0.014625475	minimum_validation_cost	0.012483893
epoch no :	3500	training cost:	0.01181775	validation cost:	0.013272421	minimum_validation_cost	0.012483893
epoch no :	4000	training cost:	0.013972165	validation cost:	0.015514897	minimum_validation_cost	0.012483893
epoch no :	4500	training cost:	0.012848759	validation cost:	0.014008677	minimum_validation_cost	0.012483893
epoch no :	5000	training cost:	0.01293095	validation cost:	0.014150724	minimum_validation_cost	0.012483893
epoch no :	5500	training cost:	0.015687615	validation cost:	0.014634018	minimum_validation_cost	0.012483893
epoch no :	6000	training cost:	0.013886566	validation cost:	0.015381509	minimum_validation_cost	0.012483893
epoch no :	6500	training cost:	0.016043557	validation cost:	0.016748872	minimum_validation_cost	0.012483893
epoch no :	7000	training cost:	0.014979528	validation cost:	0.016671315	minimum_validation_cost	0.012483893
epoch no :	7500	training cost:	0.016729917	validation cost:	0.019415256	minimum_validation_cost	0.012483893
epoch no :	8000	training cost:	0.0152445175	validation cost:	0.016355077	minimum_validation_cost	0.012483893
epoch no :	8500	training cost:	0.015557777	validation cost:	0.016310645	minimum_validation_cost	0.012483893
epoch no :	9000	training cost:	0.07586811	validation cost:	0.07899709	minimum_validation_cost	0.012483893
epoch no :	9500	training cost:	0.02142299	validation cost:	0.024720391	minimum_validation_cost	0.012483893
epoch no :	10000	training cost:	0.024426803	validation cost:	0.027593074	minimum_validation_cost	0.012483893
epoch no :	10500	training cost:	0.018516378	validation cost:	0.019000562	minimum_validation_cost	0.012483893
epoch no :	11000	training cost:	0.048320018	validation cost:	0.048580103	minimum_validation_cost	0.012483893
epoch no :	11500	training cost:	0.020851415	validation cost:	0.037529126	minimum_validation_cost	0.012483893
epoch no :	12000	training cost:	0.058065794	validation cost:	0.029545134	minimum_validation_cost	0.012483893
epoch no :	12500	training cost:	0.015809786	validation cost:	0.016998455	minimum_validation_cost	0.012483893
epoch no :	13000	training cost:	0.01982437	validation cost:	0.02025333	minimum_validation_cost	0.012483893
epoch no :	13500	training cost:	0.014694117	validation cost:	0.016001858	minimum_validation_cost	0.012483893
epoch no :	14000	training cost:	0.0147103565	validation cost:	0.015879849	minimum_validation_cost	0.012483893
epoch no :	14500	training cost:	0.071125045	validation cost:	0.035555735	minimum_validation_cost	0.012483893
epoch no :	15000	training cost:	0.023572352	validation cost:	0.016225783	minimum_validation_cost	0.012483893


Model saved in path: model/model.ckpt

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

ann root mean absolute error: 0.10504402603717251

accuracy score: 0.9348914236909488



Description on Learning curve and Accuracy:

We can observe where overfitting occurs. Overfitting actually occurs if the training loss goes under the validation loss even though the validation is still dropping. It is the sign that network is learning the patterns in the train set that are not applicable in the validation done. In a short note we can say::

Overfitting : training loss << validation loss

Underfitting : training loss >> validation loss

Just right : training loss ~ validation loss

According to this theory, for ANN 1,2 and 3 our both learning curve (validation cost and training cost) is exactly top of one another so in our case validation cost and training cost is almost same so we can say that our model is doing just the right thing. Again In validation score .11,.1081 and .1050 is impressive compared to other models.

ANN single hidden layer

```
In [281]: tf.reset_default_graph()
def weight_bais():
    global weights, biases
    weights = {
        'w1': tf.Variable(tf.random_normal([input_dim, hidden_1])),
        'out': tf.Variable(tf.random_normal([hidden_1, output_dim]))
    }
    biases = {
        'b1': tf.Variable(tf.random_normal([hidden_1])),
        'out': tf.Variable(tf.random_normal([output_dim]))
    }
```



```
In [282]: def ann_model(X_val):
# Hidden layers
layer_1 = tf.add(tf.matmul(X_val, weights['w1']), biases['b1'])
layer_1 = tf.nn.relu(layer_1)

# Output layer
layer_out = tf.add(tf.matmul(layer_1, weights['out']), biases['out'])

return layer_out
```

```
In [283]: def miscellaneous_initialization():
global model, cost , regularizer_1 , regularizer_2 ,regularizer_3, regularizer_4, optimizer , init , saver
# Model Construct
model = ann_model(X_tf)

# Mean Squared Error cost function
cost = tf.reduce_mean(tf.square(y_tf - model))

# cost = tf.square(y_tf - model)
regularizer_1 = tf.nn.l2_loss(weights['w1'])

# cost = tf.reduce_mean(cost + beta1*regularizer_1 + beta2*regularizer_2 + beta3*regularizer_3)
cost = cost + beta1*regularizer_1

# Adam optimizer will update weights and biases after each step
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)

# Initialize variables
init = tf.global_variables_initializer()

# Add ops to save and restore all the variables.
saver = tf.train.Saver()
```

ANN 4

- learning rate = .1

layer name	Neuron	value of beta for l2 regularization
1st hidden layer	16 Neuron	.1

```

In [284]: tf.reset_default_graph()
learning_rate = 0.1
num_steps = 15000
#for regularize weight matrix
beta1 = 0.1
beta2 = None
beta3 = None
beta4 = None
minimum_validation_cost = 0.01901000
hidden_1 = 16
hidden_2 = None
hidden_3 = None
hidden_4 = None

#tf graph input
X_tf = tf.placeholder("float" )
y_tf = tf.placeholder("float" )

weight_bais()
miscellaneous_initialization()
train_LC = []
val_LC = []
training_block(X_train,y_train, X_test,y_test)
prediction = Prediction_block(X_test)

test_rmse_score, test_r2_score = accuracy(y_test,prediction)

print('ann root mean absolute error: ', test_rmse_score)
print('accuracy score: ', test_r2_score )
learning_curve()

pred_str = 'ANN_lr'+str(learning_rate)+'_beta'+str(beta1)+'-'+str(beta2)+'-'+str(beta3)+'-'+str(beta4)+'_hidden'+str(hidden_1)+'-'+str(hidden_2)+'-'+str(hidden_3)+'-'+str(hidden_4)
prediction_dict[pred_str] = prediction

if submit:
    submit_prediction = Prediction_block(test_processed)
    submit_prediction_dict[pred_str] = submit_prediction
# Data Save
if save_score:
    log_df = pd.read_csv("different_training_results.csv")
    log_df = log_df.append({'learning_rate' : learning_rate, 'num_steps' : num_steps, 'beta1' : beta1, 'beta2' : beta2, 'beta3' : beta3, 'beta4' : beta4, 'hidden_1' : hidden_1 , 'hidden_2' : hidden_2, 'hidden_3' : hidden_3, 'hidden_4' : hidden_4, 'input_dim' : input_dim , 'test_rmse_score' : test_rmse_score , 'test_r2_score' : test_r2_score}, ignore_index=True)
    log_df.to_csv("different_training_results.csv", encoding='utf-8',index=False)

```

epoch no :	500	training cost:	0.70696795	validation cost:	0.70672375	minimum_validation_cost	0.01901
epoch no :	1000	training cost:	0.074884504	validation cost:	0.09517263	minimum_validation_cost	0.01901
epoch no :	1500	training cost:	0.030514441	validation cost:	0.037706994	minimum_validation_cost	0.01901
epoch no :	2000	training cost:	0.041105207	validation cost:	0.06775436	minimum_validation_cost	0.0175587
epoch no :	2500	training cost:	0.040804587	validation cost:	0.048201256	minimum_validation_cost	0.015757574
epoch no :	3000	training cost:	0.023604002	validation cost:	0.020474989	minimum_validation_cost	0.01474772
epoch no :	3500	training cost:	0.021396361	validation cost:	0.02962465	minimum_validation_cost	0.013838205
epoch no :	4000	training cost:	0.011788214	validation cost:	0.013150874	minimum_validation_cost	0.013117214
epoch no :	4500	training cost:	0.011998331	validation cost:	0.013331305	minimum_validation_cost	0.013110494
epoch no :	5000	training cost:	0.01817273	validation cost:	0.016207738	minimum_validation_cost	0.013110494
epoch no :	5500	training cost:	0.107278176	validation cost:	0.05767396	minimum_validation_cost	0.013110494
epoch no :	6000	training cost:	0.019238576	validation cost:	0.016842552	minimum_validation_cost	0.013110494
epoch no :	6500	training cost:	0.013488971	validation cost:	0.015036552	minimum_validation_cost	0.013110494
epoch no :	7000	training cost:	0.014104144	validation cost:	0.014767384	minimum_validation_cost	0.013110494
epoch no :	7500	training cost:	0.0142570175	validation cost:	0.015163885	minimum_validation_cost	0.013110494
epoch no :	8000	training cost:	0.014391301	validation cost:	0.015554772	minimum_validation_cost	0.013110494
epoch no :	8500	training cost:	0.03745286	validation cost:	0.047344368	minimum_validation_cost	0.013110494
epoch no :	9000	training cost:	0.015391441	validation cost:	0.017085707	minimum_validation_cost	0.013110494
epoch no :	9500	training cost:	0.05376236	validation cost:	0.02105463	minimum_validation_cost	0.013110494
epoch no :	10000	training cost:	0.2804591	validation cost:	0.29511937	minimum_validation_cost	0.013110494
epoch no :	10500	training cost:	0.015953306	validation cost:	0.016810209	minimum_validation_cost	0.013110494
epoch no :	11000	training cost:	0.01592241	validation cost:	0.016748453	minimum_validation_cost	0.013110494
epoch no :	11500	training cost:	0.01610453	validation cost:	0.01695231	minimum_validation_cost	0.013110494
epoch no :	12000	training cost:	0.07394206	validation cost:	0.12320263	minimum_validation_cost	0.013110494
epoch no :	12500	training cost:	0.15734634	validation cost:	0.15152527	minimum_validation_cost	0.013110494
epoch no :	13000	training cost:	0.01652914	validation cost:	0.017331216	minimum_validation_cost	0.013110494
epoch no :	13500	training cost:	0.016681496	validation cost:	0.017460056	minimum_validation_cost	0.013110494
epoch no :	14000	training cost:	0.020754531	validation cost:	0.021394152	minimum_validation_cost	0.013110494
epoch no :	14500	training cost:	0.016063426	validation cost:	0.016903006	minimum_validation_cost	0.013110494
epoch no :	15000	training cost:	0.019355612	validation cost:	0.020434875	minimum_validation_cost	0.013110494


Model saved in path: model/model.ckpt

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

ann root mean absolute error: 0.10595639845030991

accuracy score: 0.9337554952955662



ANN 5

- learning rate = .1

layer name	Neuron	value of beta for l2 regularization
1st hidden layer	4 Neuron	.1

```
In [285]: tf.reset_default_graph()
learning_rate = 0.1
num_steps = 8000
#for regularize weight matrix
beta1 = 0
beta2 = None
beta3 = None
beta4 = None

hidden_1 = 4
hidden_2 = None
hidden_3 = None
hidden_4 = None
minimum_validation_cost = 0.1701000
#tf graph input
X_tf = tf.placeholder("float" )
y_tf = tf.placeholder("float" )

weight_bais()
miscellaneous_initialization()
train_LC = []
val_LC = []
```

```

In [286]: training_block(X_train,y_train, X_test,y_test)
prediction = Prediction_block(X_test)

test_rmse_score, test_r2_score = accuracy(y_test,prediction)

print('ann root mean absolute error: ', test_rmse_score)
print('accuracy score: ', test_r2_score )
learning_curve()

pred_str = 'ANN_lr'+str(learning_rate)+'_beta'+str(beta1)+'-'+str(beta2)+'-'+str(beta3)+'-'+str(beta4)+'_hidden'+str(hidden_1)+'-'+str(hidden_2)+'-'+str(hidden_3)+'-'+str(hidden_4)
prediction_dict[pred_str] = prediction

if submit:
    submit_prediction = Prediction_block(test_processed)
    submit_prediction_dict[pred_str] = submit_prediction
# Data Save
if save_score:
    log_df = pd.read_csv("diffrent_training_results.csv")
    log_df = log_df.append({'learning_rate' : learning_rate, 'num_steps' : num_steps, 'beta1' : beta1, 'beta2' : beta2, 'beta3' : beta3, 'beta4' : beta4, 'hidden_1' : hidden_1 , 'hidden_2' : hidden_2, 'hidden_3' : hidden_3, 'hidden_4' : hidden_4, 'input_dim' : input_dim , 'test_rmse_score' : test_rmse_score , 'test_r2_score' : test_r2_score}, ignore_index=True)
    log_df.to_csv("diffrent_training_results.csv", encoding='utf-8',index=False)

```

layer name	Neuron	value of beta for l2 regularization
1st hidden layer	2 Neuron	.1

```
In [98]: tf.reset_default_graph()
learning_rate = 0.1
num_steps = 15000
#for regularize weight matrix
beta1 = 0
beta2 = None
beta3 = None
beta4 = None

hidden_1 = 2
hidden_2 = None
hidden_3 = None
hidden_4 = None
minimum_validation_cost = 0.01901000
#tf graph input
X_tf = tf.placeholder("float" )
y_tf = tf.placeholder("float" )

weight_bais()
miscellaneous_initialization()
train_LC = []
val_LC = []
```



```

In [99]: training_block(X_train,y_train, X_test,y_test)
prediction = Prediction_block(X_test)

test_rmse_score, test_r2_score = accuracy(y_test,prediction)

print('ann root mean absolute error: ', test_rmse_score)
print('accuracy score: ', test_r2_score )
learning_curve()

pred_str = 'ANN_lr'+str(learning_rate)+'_beta'+str(beta1)+'-'+str(beta2)+'-'+str(beta3)+'-'+str(beta4)+'_hidden'+str(hidden_1)+'-
'+str(hidden_2)+'-'+str(hidden_3)+'-'+str(hidden_4)
prediction_dict[pred_str] = prediction

if submit:
    submit_prediction = Prediction_block(test_processed)
    submit_prediction_dict[pred_str] = submit_prediction
# Data Save
if save_score:
    log_df = pd.read_csv("diffrent_training_results.csv")
    log_df = log_df.append({'learning_rate' : learning_rate, 'num_steps' : num_steps, 'beta1' : beta1, 'beta2' : beta2, 'beta3' :
beta3, 'beta4' : beta4, 'hidden_1' : hidden_1 , 'hidden_2' : hidden_2, 'hidden_3' : hidden_3, 'hidden_4' : hidden_4, 'input_dim'
: input_dim , 'test_rmse_score' : test_rmse_score , 'test_r2_score' : test_r2_score}, ignore_index=True)
    log_df.to_csv("diffrent_training_results.csv", encoding='utf-8',index=False)

```

epoch no :	500	training cost:	0.019526327	validation cost:	0.04785805	minimum_validation_cost	0.01901
epoch no :	1000	training cost:	0.090430334	validation cost:	0.11915722	minimum_validation_cost	0.01901
epoch no :	1500	training cost:	0.009021159	validation cost:	0.019854361	minimum_validation_cost	0.01901
epoch no :	2000	training cost:	0.008328172	validation cost:	0.018638173	minimum_validation_cost	0.017471917
epoch no :	2500	training cost:	0.007971583	validation cost:	0.018678952	minimum_validation_cost	0.017459333
epoch no :	3000	training cost:	0.0077294353	validation cost:	0.018802863	minimum_validation_cost	0.017459333
epoch no :	3500	training cost:	0.0075098113	validation cost:	0.018267034	minimum_validation_cost	0.017459333
epoch no :	4000	training cost:	0.015966242	validation cost:	0.032293662	minimum_validation_cost	0.017459333
epoch no :	4500	training cost:	0.0074046743	validation cost:	0.017824942	minimum_validation_cost	0.017459333
epoch no :	5000	training cost:	0.00737159	validation cost:	0.017943176	minimum_validation_cost	0.017459333
epoch no :	5500	training cost:	0.009448293	validation cost:	0.018865194	minimum_validation_cost	0.017379621
epoch no :	6000	training cost:	0.007176871	validation cost:	0.017734537	minimum_validation_cost	0.017366491
epoch no :	6500	training cost:	0.17015442	validation cost:	0.19888674	minimum_validation_cost	0.017291883
epoch no :	7000	training cost:	0.048054833	validation cost:	0.07152697	minimum_validation_cost	0.017291883
epoch no :	7500	training cost:	0.007178404	validation cost:	0.017213237	minimum_validation_cost	0.017266192
epoch no :	8000	training cost:	0.0071287826	validation cost:	0.017250806	minimum_validation_cost	0.017099544
epoch no :	8500	training cost:	0.078824885	validation cost:	0.047651436	minimum_validation_cost	0.016776746
epoch no :	9000	training cost:	0.007192865	validation cost:	0.017091228	minimum_validation_cost	0.016776746
epoch no :	9500	training cost:	0.0071183126	validation cost:	0.016838994	minimum_validation_cost	0.016650783
epoch no :	10000	training cost:	0.020033514	validation cost:	0.03116488	minimum_validation_cost	0.016502209
epoch no :	10500	training cost:	0.023244442	validation cost:	0.024413731	minimum_validation_cost	0.01633817
epoch no :	11000	training cost:	0.013847738	validation cost:	0.027082304	minimum_validation_cost	0.016238058
epoch no :	11500	training cost:	0.007122549	validation cost:	0.016189996	minimum_validation_cost	0.015966171
epoch no :	12000	training cost:	0.007111148	validation cost:	0.0159556	minimum_validation_cost	0.015759723
epoch no :	12500	training cost:	0.007110614	validation cost:	0.01570639	minimum_validation_cost	0.01548126
epoch no :	13000	training cost:	0.0071250037	validation cost:	0.0154177975	minimum_validation_cost	0.0152186025
epoch no :	13500	training cost:	0.0071092476	validation cost:	0.015325032	minimum_validation_cost	0.015116011
epoch no :	14000	training cost:	0.007109693	validation cost:	0.015189837	minimum_validation_cost	0.014980882
epoch no :	14500	training cost:	0.007109804	validation cost:	0.015131356	minimum_validation_cost	0.014810974
epoch no :	15000	training cost:	0.007110104	validation cost:	0.014890727	minimum_validation_cost	0.014607174


Model saved in path: model/model.ckpt

INFO:tensorflow:Restoring parameters from model/model.ckpt

Model restored.

ann root mean absolute error: 0.12086015230808778

accuracy score: 0.9138090324301579



Description on Learning curve and Accuracy:

We can observe where overfitting occurs. Overfitting actually occurs if the training loss goes under the validation loss even though the validation is still dropping. It is the sign that network is learning the patterns in the train set that are not applicable in the validation done. In a short note we can say::

Overfitting : training loss << validation loss

Underfitting : training loss >> validation loss

Just right : training loss ~ validation loss

According to this theory, for ANN 4 our both learning curve (validation cost and training cost) is exactly top of one another so in our case validation cost and training cost is almost same so we can say that our model is doing just the right thing. Again In validation score .1059 is impressive compared to other models.

But for ANN 5 and 6 training loss << validation loss so we can say that this two model overfit data due to lower amount of neuron but ANN 4 have just the right amount of neuron thats why with similar parameter this overfit occered.

Hyperparameeter tuning

Few of my hyperparameeter tuning is shown in the following block. In this data if a hidden layer value is 0 then it means that the hidden layer is turned off. For example if hidden_3 = 0 then that means hidden layer 3 is removed from the model and the model have only 2 hidden layer. And all the score is done on a validation set which is not seen by the model while training. For most of the case it was a 80-20 split. In the following results I didnt kept any cross validation results but I have used diffrent seed while splitting data due to diffrent seed sometimes good hyperparameeter also provided so so accuracy.

```
In [255]: log_df = pd.read_csv("different_training_results.csv")
# print(log_df.to_string())
pd.set_option('display.max_rows', None)
log_df
```

Out[255]:

	learning_rate	num_steps	beta1	beta2	beta3	hidden_1	hidden_2	hidden_3	input_dim	test_rmse_score	test_r2_score	hidden_4	beta4
0	0.050	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.128456	8.949361e-01	NaN	NaN
1	0.040	2500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.250470	6.005558e-01	NaN	NaN
2	0.050	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.152580	8.517686e-01	NaN	NaN
3	0.050	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.143409	8.690530e-01	NaN	NaN
4	0.050	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.127356	8.967284e-01	NaN	NaN
5	0.050	7900.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.126758	8.976948e-01	NaN	NaN
6	0.050	1500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.162495	8.318785e-01	NaN	NaN
7	0.050	1500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.177628	7.991050e-01	NaN	NaN
8	0.050	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.139909	8.753655e-01	NaN	NaN
9	0.050	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.143775	8.683835e-01	NaN	NaN
10	0.100	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.138477	8.779036e-01	NaN	NaN
11	0.100	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.138477	8.779036e-01	NaN	NaN
12	0.010	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.154219	8.485668e-01	NaN	NaN
13	0.010	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.154219	8.485668e-01	NaN	NaN
14	0.010	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.661086	-1.782662e+00	NaN	NaN
15	0.100	3500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.131423	8.900259e-01	NaN	NaN
16	0.100	2800.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.389771	-5.111744e-03	NaN	NaN
17	0.100	2500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.390269	-7.679426e-03	NaN	NaN
18	0.100	8500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.102641	9.302995e-01	NaN	NaN
19	0.050	8500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.135519	8.830642e-01	NaN	NaN
20	0.100	2500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.304550	4.094407e-01	NaN	NaN
21	0.100	9500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.124793	9.008422e-01	NaN	NaN
22	0.001	9500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	14.246912	-1.291369e+03	NaN	NaN
23	0.005	9500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.710445	-2.213707e+00	NaN	NaN
24	0.100	7500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.129652	8.929701e-01	NaN	NaN
25	0.100	12500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.142223	8.712101e-01	NaN	NaN
26	0.100	11500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.126000	8.989146e-01	NaN	NaN
27	0.050	11500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.124008	9.020864e-01	NaN	NaN
28	0.050	11000.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.147886	8.607480e-01	NaN	NaN
29	0.050	11000.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.105610	9.262080e-01	NaN	NaN
30	0.050	11000.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.131099	8.905686e-01	NaN	NaN
31	0.010	11000.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	2.365656	-3.463267e+01	NaN	NaN

	learning_rate	num_steps	beta1	beta2	beta3	hidden_1	hidden_2	hidden_3	input_dim	test_rmse_score	test_r2_score	hidden_4	beta4
32	0.100	11000.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.135869	8.824604e-01	NaN	NaN
33	0.100	11000.0	0.005	0.0050	0.000000	16.0	8.0	4.0	403.0	0.132422	8.883490e-01	NaN	NaN
34	0.050	13000.0	0.005	0.0050	0.000000	16.0	8.0	4.0	403.0	0.126124	8.987160e-01	NaN	NaN
35	0.100	23000.0	0.005	0.0050	0.000000	16.0	8.0	4.0	403.0	0.123945	9.021856e-01	NaN	NaN
36	0.100	23000.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.321211	3.430598e-01	NaN	NaN
37	0.050	23000.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.126705	8.977802e-01	NaN	NaN
38	0.050	23000.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.126705	8.977802e-01	NaN	NaN
39	0.050	23000.0	0.000	0.0000	0.000000	16.0	8.0	4.0	403.0	0.173484	8.083707e-01	NaN	NaN
40	0.050	23000.0	0.000	0.0000	0.000000	16.0	8.0	4.0	403.0	0.396362	-2.986747e-04	NaN	NaN
41	0.100	23000.0	0.000	0.0000	0.000000	16.0	8.0	4.0	403.0	0.396368	-3.297183e-04	NaN	NaN
42	0.100	17000.0	0.000	0.0000	0.000000	16.0	8.0	4.0	403.0	0.396362	-2.981759e-04	NaN	NaN
43	0.001	3000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	2.654960	-4.388086e+01	NaN	NaN
44	0.001	3000.0	0.050	0.0000	0.000000	16.0	8.0	4.0	403.0	7.177147	-3.269811e+02	NaN	NaN
45	0.100	3000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.124692	9.010024e-01	NaN	NaN
46	0.100	3000.0	0.010	0.0000	0.000000	16.0	8.0	4.0	403.0	0.126720	8.977569e-01	NaN	NaN
47	0.100	13000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.396362	-2.986747e-04	NaN	NaN
48	0.100	2500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.136049	8.821482e-01	NaN	NaN
49	0.100	3000.0	0.100	0.1000	0.000000	16.0	8.0	4.0	403.0	0.127360	8.967204e-01	NaN	NaN
50	0.100	4000.0	0.100	0.1000	0.000000	16.0	8.0	4.0	403.0	0.260455	5.680732e-01	NaN	NaN
51	0.100	3500.0	0.100	0.0100	0.000000	16.0	8.0	4.0	403.0	0.146195	8.639153e-01	NaN	NaN
52	0.100	3500.0	0.100	0.0010	0.000000	16.0	8.0	4.0	403.0	0.396370	-3.376305e-04	NaN	NaN
53	0.100	3500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.113379	9.149525e-01	NaN	NaN
54	0.100	3500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.117934	9.079823e-01	NaN	NaN
55	0.100	8500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.111485	9.177711e-01	NaN	NaN
56	0.100	7600.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.106919	9.243687e-01	NaN	NaN
57	0.050	7600.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.119165	9.060501e-01	NaN	NaN
58	0.100	7600.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.376772	6.081402e-02	NaN	NaN
59	0.100	3500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.107788	9.231341e-01	NaN	NaN
60	0.050	8500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.122025	9.014877e-01	NaN	NaN
61	0.100	7600.0	0.005	0.0050	0.000000	16.0	8.0	4.0	403.0	0.389771	-5.109640e-03	NaN	NaN
62	0.100	7600.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.109705	9.203748e-01	NaN	NaN
63	0.100	7600.0	0.100	0.0050	0.005000	200.0	100.0	30.0	403.0	0.389676	-4.622793e-03	NaN	NaN

	learning_rate	num_steps	beta1	beta2	beta3	hidden_1	hidden_2	hidden_3	input_dim	test_rmse_score	test_r2_score	hidden_4	beta4
64	0.100	9600.0	0.000	0.0000	0.000000	200.0	100.0	30.0	403.0	0.389825	-5.391521e-03	NaN	NaN
65	0.100	3600.0	0.100	0.0000	0.000000	200.0	100.0	30.0	403.0	1.764792	-1.960547e+01	NaN	NaN
66	0.100	7600.0	0.100	0.0000	0.000000	200.0	100.0	30.0	403.0	0.389774	-5.127892e-03	NaN	NaN
67	0.100	17600.0	0.100	0.0000	0.000000	200.0	100.0	30.0	403.0	0.389772	-5.119113e-03	NaN	NaN
68	0.100	15600.0	0.100	0.0000	0.000000	200.0	100.0	30.0	403.0	0.149463	8.522031e-01	NaN	NaN
69	0.100	15600.0	0.100	0.0000	0.000000	32.0	16.0	8.0	403.0	0.389750	-5.006016e-03	NaN	NaN
70	0.100	3600.0	0.100	0.0000	0.000000	1.0	0.0	0.0	403.0	0.104633	9.275680e-01	NaN	NaN
71	0.100	7500.0	0.100	0.0000	0.000000	1.0	0.0	0.0	403.0	0.105433	9.264551e-01	NaN	NaN
72	0.100	7500.0	0.010	0.0000	0.000000	1.0	0.0	0.0	403.0	0.104463	9.278026e-01	NaN	NaN
73	0.100	6000.0	0.010	0.0000	0.000000	1.0	0.0	0.0	403.0	0.158149	8.345259e-01	NaN	NaN
74	0.100	86000.0	0.010	0.0000	0.000000	1.0	0.0	0.0	403.0	0.389882	-5.682449e-03	NaN	NaN
75	0.100	8600.0	0.010	0.0000	0.000000	1.0	0.0	0.0	403.0	0.389771	-5.109640e-03	NaN	NaN
76	0.100	8600.0	0.100	0.0000	0.000000	1.0	0.0	0.0	403.0	0.111076	9.183720e-01	NaN	NaN
77	0.100	3600.0	0.100	0.0000	0.000000	4.0	0.0	0.0	403.0	0.105270	9.266836e-01	NaN	NaN
78	0.100	3600.0	0.100	0.0000	0.000000	1.0	0.0	0.0	403.0	0.104898	9.271996e-01	NaN	NaN
79	0.100	3600.0	0.100	0.0000	0.000000	32.0	0.0	0.0	403.0	0.198056	7.404795e-01	NaN	NaN
80	0.100	3600.0	0.100	0.0000	0.000000	16.0	0.0	0.0	403.0	0.607723	-1.443464e+00	NaN	NaN
81	0.100	3600.0	0.100	0.0000	0.000000	16.0	0.0	0.0	403.0	0.106043	9.256030e-01	NaN	NaN
82	0.100	7600.0	0.100	0.0000	0.000000	1.0	0.0	0.0	403.0	0.105914	9.257826e-01	NaN	NaN
83	0.100	7600.0	0.100	0.0000	0.000000	16.0	0.0	0.0	403.0	0.107050	9.241823e-01	NaN	NaN
84	0.100	3900.0	0.100	0.0000	0.000000	16.0	0.0	0.0	403.0	0.107679	9.232893e-01	NaN	NaN
85	0.100	2000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.157928	8.411961e-01	NaN	NaN
86	0.100	7600.0	0.100	0.0050	0.005000	16.0	8.0	4.0	403.0	0.138570	8.777409e-01	NaN	NaN
87	0.100	8600.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.133445	8.866164e-01	NaN	NaN
88	0.100	8600.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.129291	8.935660e-01	NaN	NaN
89	0.100	3600.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.119097	9.096871e-01	NaN	NaN
90	0.100	3500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.512748	-6.739936e-01	NaN	NaN
91	0.100	9600.0	0.100	0.0050	0.005000	16.0	8.0	4.0	403.0	0.124896	9.006785e-01	NaN	NaN
92	0.100	9600.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.205646	7.307324e-01	NaN	NaN
93	0.100	19600.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.132143	8.888180e-01	NaN	NaN
94	0.100	19600.0	0.100	0.0100	0.001000	16.0	8.0	4.0	403.0	0.235628	6.464929e-01	NaN	NaN
95	0.100	19600.0	0.100	0.0005	0.000005	16.0	8.0	4.0	403.0	0.128857	8.942789e-01	NaN	NaN

	learning_rate	num_steps	beta1	beta2	beta3	hidden_1	hidden_2	hidden_3	input_dim	test_rmse_score	test_r2_score	hidden_4	beta4
96	0.100	19600.0	0.000	0.0000	0.000000	16.0	8.0	4.0	403.0	0.396391	-4.451594e-04	NaN	NaN
97	0.100	29600.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.148700	8.592123e-01	NaN	NaN
98	0.100	4000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.172257	8.110704e-01	NaN	NaN
99	0.100	1750.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.396362	-2.980097e-04	NaN	NaN
100	0.100	3600.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.127960	8.957456e-01	NaN	NaN
101	0.100	4000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.122812	9.039655e-01	NaN	NaN
102	0.100	7600.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.127266	8.968736e-01	NaN	NaN
103	0.100	5500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.127831	8.959562e-01	NaN	NaN
104	0.100	7600.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	399.635052	-1.016885e+06	NaN	NaN
105	0.100	17600.0	0.100	0.0005	0.000005	16.0	8.0	4.0	403.0	0.122934	9.037754e-01	NaN	NaN
106	0.100	17100.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.123996	9.021051e-01	NaN	NaN
107	0.100	19500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.122128	9.050330e-01	NaN	NaN
108	0.100	29500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.121803	9.055369e-01	NaN	NaN
109	0.100	49500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.120712	9.072223e-01	NaN	NaN
110	0.100	9500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.104178	9.281964e-01	NaN	NaN
111	0.100	19500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.089601	9.468848e-01	NaN	NaN
112	0.100	19500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.118142	9.076568e-01	NaN	NaN
113	0.100	19500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.118142	9.076568e-01	NaN	NaN
114	0.100	9500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.090853	9.453897e-01	NaN	NaN
115	0.100	9500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.090915	9.453158e-01	NaN	NaN
116	0.100	19500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.091638	9.444420e-01	NaN	NaN
117	0.050	39500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.076999	9.607748e-01	NaN	NaN
118	0.100	49500.0	0.005	0.0050	0.005000	16.0	8.0	4.0	403.0	0.073579	9.641818e-01	NaN	NaN
119	0.100	49500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.087605	9.492243e-01	NaN	NaN
120	0.100	9500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.159089	8.325543e-01	NaN	NaN
121	0.100	29500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.093281	9.424324e-01	NaN	NaN
122	0.100	29500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.127969	8.689350e-01	NaN	NaN
123	0.100	9500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.117637	8.892446e-01	NaN	NaN
124	0.100	9500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.104068	9.283477e-01	NaN	NaN
125	0.100	9500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.103677	9.288849e-01	NaN	NaN
126	0.100	9500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.103494	9.291363e-01	NaN	NaN
127	0.100	9500.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.103775	9.364546e-01	NaN	NaN


	learning_rate	num_steps	beta1	beta2	beta3	hidden_1	hidden_2	hidden_3	input_dim	test_rmse_score	test_r2_score	hidden_4	beta4
128	0.100	26000.0	0.100	0.0000	0.000000	200.0	100.0	50.0	403.0	0.112221	9.256909e-01	NaN	NaN
129	0.100	26000.0	0.100	0.0000	0.000000	200.0	100.0	50.0	403.0	0.111169	9.270777e-01	12.0	NaN
130	0.100	26000.0	0.100	0.0000	0.000000	200.0	100.0	50.0	403.0	0.111169	9.270777e-01	12.0	NaN
131	0.100	46000.0	0.100	0.0000	0.000000	200.0	100.0	50.0	403.0	0.411770	-4.727709e-04	12.0	NaN
132	0.100	26000.0	0.000	0.0000	0.000000	200.0	96.0	32.0	403.0	0.411734	-2.970530e-04	4.0	NaN
133	0.100	6000.0	0.000	0.0000	0.000000	200.0	96.0	32.0	403.0	0.412326	-3.176449e-03	4.0	0.0
134	0.100	26000.0	0.100	0.0000	0.000000	200.0	96.0	32.0	403.0	0.411762	-4.307332e-04	4.0	0.0
135	0.100	26000.0	0.100	0.0000	0.000000	32.0	16.0	8.0	403.0	0.411745	-3.482792e-04	4.0	0.0
136	0.100	46000.0	0.100	0.1000	0.000000	200.0	100.0	50.0	403.0	0.411750	-3.748811e-04	25.0	0.0
137	0.100	46000.0	0.100	0.1000	0.000000	200.0	100.0	50.0	403.0	0.411750	-3.748811e-04	25.0	0.0
138	0.100	26000.0	0.100	0.0000	0.000000	200.0	100.0	50.0	403.0	0.411739	-3.223233e-04	12.0	0.0
139	0.050	26000.0	0.100	0.0000	0.000000	200.0	100.0	50.0	403.0	0.412324	-3.166794e-03	12.0	0.0
140	0.100	15000.0	0.100	NaN	NaN	16.0	NaN	NaN	403.0	0.115072	9.218665e-01	NaN	NaN
141	0.100	15000.0	0.000	NaN	NaN	1.0	NaN	NaN	403.0	0.116185	9.203481e-01	NaN	NaN
142	0.100	15000.0	0.000	NaN	NaN	2.0	NaN	NaN	403.0	0.150994	8.654721e-01	NaN	NaN
143	0.100	15000.0	NaN	NaN	NaN	NaN	NaN	NaN	403.0	0.162648	8.439027e-01	NaN	NaN
144	0.100	35000.0	NaN	NaN	NaN	NaN	NaN	NaN	403.0	0.147050	8.724072e-01	NaN	NaN
145	0.100	8000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.184813	7.984608e-01	NaN	NaN
146	0.100	20000.0	0.100	0.0500	0.000000	76.0	48.0	32.0	403.0	0.111441	9.267196e-01	16.0	0.0
147	0.100	25000.0	0.100	0.0500	0.000000	128.0	64.0	32.0	403.0	0.412339	-3.237055e-03	16.0	0.0
148	0.100	8000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.111108	9.271576e-01	NaN	NaN
149	0.100	25000.0	0.100	0.0500	0.000000	128.0	64.0	32.0	403.0	0.150787	8.658403e-01	16.0	0.0
150	0.050	35000.0	0.100	0.0000	0.000000	128.0	64.0	32.0	403.0	0.412324	-3.166794e-03	16.0	0.0
151	0.100	25000.0	0.100	0.0000	0.000000	256.0	128.0	32.0	403.0	0.412324	-3.166794e-03	8.0	0.0
152	0.050	25000.0	0.100	0.0000	0.000000	256.0	128.0	32.0	403.0	0.412324	-3.167316e-03	8.0	0.0
153	0.100	15000.0	0.100	0.0000	0.000000	4.0	16.0	16.0	403.0	0.412324	-3.167055e-03	4.0	0.0
154	0.100	35000.0	0.100	0.0000	0.000000	256.0	128.0	64.0	403.0	0.412325	-3.170185e-03	32.0	0.0
155	0.100	25000.0	0.100	0.1000	0.000000	256.0	128.0	32.0	403.0	0.123334	9.102441e-01	8.0	0.0
156	0.050	25000.0	0.100	0.0000	0.000000	256.0	128.0	64.0	403.0	0.412324	-3.167055e-03	8.0	0.0
157	0.050	25000.0	0.100	0.0000	0.000000	128.0	64.0	16.0	403.0	0.412319	-3.141554e-03	4.0	0.0
158	0.100	35000.0	0.100	0.0000	0.000000	16.0	32.0	48.0	403.0	0.416205	-2.213767e-02	76.0	0.0
159	0.100	15000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.412324	-3.166794e-03	2.0	0.0

	learning_rate	num_steps	beta1	beta2	beta3	hidden_1	hidden_2	hidden_3	input_dim	test_rmse_score	test_r2_score	hidden_4	beta4
160	0.050	25000.0	0.100	0.0500	0.000000	128.0	64.0	16.0	403.0	0.412324	-3.167837e-03	4.0	0.0
161	0.050	25000.0	0.100	0.0000	0.000000	190.0	90.0	30.0	403.0	0.412324	-3.165491e-03	3.0	0.0
162	0.100	8000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.412324	-3.166794e-03	2.0	0.0
163	0.100	6000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.412324	-3.165230e-03	NaN	NaN
164	0.100	6000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.105436	9.344041e-01	NaN	NaN
165	0.100	20000.0	0.100	0.0500	0.000000	76.0	48.0	32.0	403.0	0.108646	9.303500e-01	16.0	0.0
166	0.100	6000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.103792	9.364344e-01	NaN	NaN
167	0.100	20000.0	0.100	0.0500	0.000000	76.0	48.0	32.0	403.0	0.106511	9.330597e-01	16.0	0.0
168	0.100	25000.0	0.100	0.0500	0.000000	128.0	64.0	32.0	403.0	0.411699	-1.253480e-04	16.0	0.0
169	0.100	25000.0	0.100	0.1000	0.000000	256.0	128.0	32.0	403.0	0.411756	-4.026479e-04	8.0	0.0
170	0.050	25000.0	0.100	0.0500	0.000000	256.0	128.0	16.0	403.0	0.112563	9.252368e-01	4.0	0.0
171	0.050	25000.0	0.100	0.0000	0.000000	190.0	90.0	30.0	403.0	0.411974	-1.465819e-03	3.0	0.0
172	0.100	8000.0	0.100	0.0000	0.000000	16.0	8.0	4.0	403.0	0.103529	9.367565e-01	2.0	0.0
173	0.100	15000.0	0.100	NaN	NaN	16.0	NaN	NaN	403.0	0.106009	9.336896e-01	NaN	NaN
174	0.100	15000.0	0.000	NaN	NaN	1.0	NaN	NaN	403.0	0.120667	9.140850e-01	NaN	NaN
175	0.100	15000.0	0.000	NaN	NaN	2.0	NaN	NaN	403.0	0.116812	9.194859e-01	NaN	NaN
176	0.100	35000.0	NaN	NaN	NaN	NaN	NaN	NaN	403.0	0.138819	8.862908e-01	NaN	NaN
177	0.050	25000.0	0.100	0.0100	0.000000	180.0	90.0	30.0	403.0	0.411965	-1.419538e-03	6.0	0.0
178	0.100	20000.0	0.100	0.0500	0.000000	76.0	48.0	32.0	403.0	0.110364	9.281301e-01	16.0	0.0


Observation and discovery :

- In the above parameter we can see that index 44 shows that for .001 learning parameter the model does not predict anything so I have changed it slowly and finally What I have found that learning parameter .1 and .05 provides the best results.
- Beta1, Beta2, Beta3, Beta4 represents the regularization parameter for hidden layer 1 ,2 ,3 and 4. Sometimes in the above table we can see that hidden layer 2,3,4 is 0 or NaN but there is some value for beta 2,3,4 that means the layer is actually off so those values actually means nothing.
- For 3 layer model when beta1, beta2, beta3 is .005, model shows significant amount of improvement while learning rate is .1 or .05 . But when learning rate is .1 and beta1=.1 , beta2=0, beta3=0 then the model performs even better most of the time and it also takes less epochs to train for the best validation accuracy
- From index 63 to 69 I have tried to use 200 , 100 , 30 neurons because the data have 403 features and its a common practice to use half amount of the neuron in the first hidden layer and this strategy does not work good enough but with my selected parameter it improved a little bit. I have used 16-8-4 combination of neuron because of this common practice. for our case 16 neuron in the first layer provided better accuracy and adding 8 and 4 in the next 2 layer improved the stability of the model and now it gives good validation accuracy after 2000 epoch and the best validation accuracy remains between the epoch range of 2000-2500 , 3300-3600 or 5000-5400 .
- From index 70 to 78 we can see that single neuron with single hidden layer performs well according to the plan stated in the target section. Then I have increased neurons and the learning curve for them is in the following block. Where y axis shows rmse and x axis shows i and i*50 represents the epoch no. Again blue curve is for training accuracy and green for validation accuracy
- In the table index 169 and 155 the model is exactly same with same parameter but one of them is providing .123 and other is providing .41 and that shows how inconsistent model become when we increase the neuron of the first hidden layer.
- We can see that even after adding another layer ANN does not perform well when we are increasing neurons in the first layer. The reason behind it is that this type of regression problem usually do well with logistic regression. By increasing neurons we cant do much improvement and all we need to do is properly regularize small amount of neurons so that they can perform well.


For single neuron learning rate



For 4 neuron learning curve



For 16 neuron learning curve



For 32 neuron learning curve

Ensemble

I am using bagging method for this section. Usually in this technique we add different models results and average them. But instead of averaging I am taking different fraction from different models result. Finally making sure that it sums up to 1.

I have tried different combinations of ensemble learning to improve performance. Kaggle has a certain limitation on uploading submission files. So what I have tried is that before submitting it to kaggle, I have made 80-20 split. I made prediction on the 20% data. Then I have tried ensemble learning so that before submission I can confirm which combination might work well.

```
In [304]: if submit :
            pred_df = pd.read_csv("diffrent_pred_results.csv")
        else:
            pred_df = pd.read_csv("pred_results.csv")

        if not submit:
            pd.set_option('display.max_colwidth', -1)
            pred_df = pd.DataFrame(prediction_dict)
            pred_df.to_csv("pred_results.csv", encoding='utf-8',index=False)

        else:
            pd.set_option('display.max_colwidth', -1)
            pred_df = pd.DataFrame(submit_prediction_dict)
            pred_df.to_csv("diffrent_pred_results.csv", encoding='utf-8',index=False)

pd.DataFrame(pred_df.columns)
```

Out[304]:

	0
0	Random Forest Regressor
1	DecisionTree
2	Xgboost
3	Lasso
4	ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None
5	ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16
6	ANN_lr0.05_beta0.005-0.1-0.05-0.0_hidden8-32-16-8
7	ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2
8	ANN_lr0.1_beta0.1-None-None-None_hidden16-None-None-None
9	ANN_lr0.1_beta0-None-None-None_hidden4-None-None-None
10	ANN_lr0.1_beta0-None-None-None_hidden2-None-None-None

Naming explanation of above table

	Name	learning rate	beta1	beta 2	beta 3	beta 4	hidden layer 1	hidden layer 2	hidden layer 3	hidden layer 4
	ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None	0.1	0.1	0.0	0.0	None	16	8	4	None
	ANN_lr0.05_beta0.005-0.1-0.05-0.0_hidden8-32-16-8	0.05	0.005	0.1	0.05	0	8	32	16	6
	ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2	.05	0.1	0.0	0.0	0.0	16	8	4	2
	ANN_lr0.1_beta0-None-None-None_hidden2-None-None-None	0.1	0	None	None	None	2	None	None	None
	ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16	0.1	.1	0.05	0.0	0.0	76	48	32	16
	ANN_lr0.1_beta0.1-None-None-None_hidden16-None-None-None	0.1	0.1	None	None	None	16	None	None	None
	ANN_lr0.1_beta0-None-None-None_hidden4-None-None-None	0.1	0	None	None	None	4	None	None	None

Ensemble Combination 1

```
In [317]: # pred_df[pred_df.columns[[1,3,5]]] * [1,2,30]
```

```
print('Using ', pred_df.columns[[4,3,2]].values)
```

```
Using      ['ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None' 'Lasso' 'Xgboost']
```

```
In [306]: prediction = pred_df[pred_df.columns[[4,3,2]]] * [.4,.2,.4]
prediction = prediction.sum(axis = 1)
```

```
if not submit:
    test_rmse_score, test_r2_score = accuracy(y_test, prediction)

    print('ann root mean absolute error: ', test_rmse_score)
    print('accuracy score: ', test_r2_score )
```

```
ann root mean absolute error: 0.1017685955982393
accuracy score: 0.9388884858512507
```


Kaggle score

[output.csv](#)

0.12231



21 hours ago by [navid](#)

Using ['ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None' 'Lasso' 'Xgboost'] *
[.4,.2,.4]

[output.csv](#)

0.12297



21 hours ago by [navid](#)

Using ['ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None' 'Lasso' 'Xgboost'] *
[.4,.2,.4]

```
In [307]: prediction = pred_df[pred_df.columns[[4,3,2]]] * [.4,.3,.3]
prediction = prediction.sum(axis = 1)

if not submit:
    test_rmse_score, test_r2_score = accuracy(y_test, prediction)

    print('ann root mean absolute error: ', test_rmse_score)
    print('accuracy score: ', test_r2_score )
```

```
ann root mean absolute error: 0.1017812588346656
accuracy score: 0.9388732764893758
```

Ensemble Combination 2

```
In [308]: # pred_df[pred_df.columns[[1,3,5]]] * [1,2,30]

print('Using ' , pred_df.columns[[2,4,5,6,7]].values)
```

```
Using ['Xgboost' 'ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None'
'ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16'
'ANN_lr0.05_beta0.005-0.1-0.05-0.0_hidden8-32-16-8'
'ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2']
```

```
In [320]: prediction = pred_df[pred_df.columns[[2,4,5,6,7]]] * [.25,.2,.2 ,.15 , .2]
prediction = prediction.sum(axis = 1)

if not submit:
    test_rmse_score, test_r2_score = accuracy(y_test, prediction)

    print('ann root mean absolute error: ', test_rmse_score)
    print('accuracy score: ', test_r2_score )
```

ann root mean absolute error: 0.1027261129768247
accuracy score: 0.9377331075112767

Kaggle score

output.csv

21 hours ago by [navid](#)

Using ['Xgboost' 'ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None'
'ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16' 'ANN_lr0.05_beta0.005-0.1-0.05-
0.0_hidden8-32-16-8' 'ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2'] * [.25,.2,.2 ,.15 , .2]

0.12319



Ensemble Combination 3

```
In [310]: print('Using ', pred_df.columns[[0,4,5,6,7]].values)
```

Using ['Random Forest Regressor'
'ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None'
'ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16'
'ANN_lr0.05_beta0.005-0.1-0.05-0.0_hidden8-32-16-8'
'ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2']

```
In [311]: prediction = pred_df[pred_df.columns[[0,4,5,6,7]]] * [.25,.2,.2 ,.15 , .2]
prediction = prediction.sum(axis = 1)

if not submit:
    test_rmse_score, test_r2_score = accuracy(y_test, prediction)

    print('ann root mean absolute error: ', test_rmse_score)
    print('accuracy score: ', test_r2_score )
```

ann root mean absolute error: 0.10392040423793271
accuracy score: 0.9362768645933334

Kaggle score

output.csv

21 hours ago by [navid](#)

0.12319



Using ['Random Forest Regressor' 'ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None' 'ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16' 'ANN_lr0.05_beta0.005-0.1-0.05-0.0_hidden8-32-16-8' 'ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2'] * [.25,.2,.2,.15,.2]

Ensemble Combination 4

```
In [316]: print('Using ', pred_df.columns[[4,5,6,7,0,2,3]].values)
```

```
Using ['ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None'
'ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16'
'ANN_lr0.05_beta0.005-0.1-0.05-0.0_hidden8-32-16-8'
'ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2' 'Random Forest Regressor'
'Xgboost' 'Lasso']
```

```
In [315]: prediction = pred_df[pred_df.columns[[4,5,6,7,0,2,3]]] * [.15,.1,.1,.05,.0,.2,.4]
prediction = prediction.sum(axis = 1)
```

```
if not submit:
    test_rmse_score, test_r2_score = accuracy(y_test, prediction)

    print('ann root mean absolute error: ', test_rmse_score)
    print('accuracy score: ', test_r2_score )
```

```
ann root mean absolute error: 0.10192734103655354
accuracy score: 0.9386976855299631
```

Kaggle score

output.csv

18 hours ago by [Navid](#)

0.12192



Using ['ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None' 'ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16' 'ANN_lr0.05_beta0.005-0.1-0.05-0.0_hidden8-32-16-8' 'ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2' 'Random Forest Regressor' 'Xgboost' 'Lasso'] * [.15,.1,.1,.05,.0,.2,.4]

Ensemble combination 4 provides the best score which is 0.12192. Currently combination 4 is showing that rmse value is .1019 and there is a better value present in combination 1 which is 0.1017. The reason behind the difference is ANN does not perform exactly same each time. That means if I currently submit with Combination 1 I might get better results than Combination 4 .

In Combination 4 used models with their parameters:

	Name	learning rate	beta1	beta 2	beta 3	beta 4	hidden layer 1	hidden layer 2	hidden layer 3	hidden layer 4	Fraction taken
	ANN_base_lr0.1_beta0.1-0.0-0.0-None_hidden16-8-4-None	0.1	0.1	0.0	0.0	None	16	8	4	None	.15
	ANN_lr0.05_beta0.005-0.1-0.05-0.0_hidden8-32-16-8	0.05	0.005	0.1	0.05	0	8	32	16	6	.1
	ANN_lr0.05_beta0.1-0.0-0.0-0.0_hidden16-8-4-2	.05	0.1	0.0	0.0	0.0	16	8	4	2	.05
	ANN_lr0.1_beta0.1-0.05-0.0-0.0_hidden76-48-32-16	0.1	.1	0.05	0.0	0.0	76	48	32	16	.1
	Xgboost	0.05	Not applicable	Not applicable	Not applicable	Not applicable	Not applicable	Not applicable	NonNot applicable	Not applicable	.2
	Lasso	alpha = 5e-4	Not applicable	Not applicable	Not applicable	Not applicable	Not applicable	Not applicable	NonNot applicable	Not applicable	.4

Obesrvation

In the learning curve graph if the minimum of training and validation is close to each other then its good to use that model. Again if training minimum and validation minimum is no where near each other then using them does not help much most of the case. When both of them are close we can use the epoch no of the train_min loss as val_min loss epoch no and then we can train over all the dataset without depending on the epoch number. The model does not give same result in same epoch every time. This is the main reason behind removing the epoch dependency.

Prepare Submission File

To use this section please uncomment the last line of split data section and comment accuracy section.

```

In [51]: use_ensemble = True
         #if want to use given test data
         if submit:
             X_test = test_processed
             if not use_ensemble:
                 with tf.Session() as sess:
                     # Restore variables from disk.
                     saver.restore(sess, "model/model.ckpt")
                     print("Model restored.")

                     # Check the values of the variables
                     pred = sess.run(model, feed_dict={X_tf: X_test})
                     prediction = pred.squeeze()

             prediction = np.exp(prediction.values)

         pred_out_df = pd.DataFrame(prediction, index=test["Id"], columns=["SalePrice"])
         pred_out_df.to_csv('output.csv', header=True, index_label='Id')

```

Conclusion & Kaggle score Discussion:

My target of this report was to improve ANN model and show how well it can perform with ANN model. In the beginning of the report I have build a ANN model that performs better than any other single ANN model. I have performed cross validation on that model and that model scored 0.12324 in kaggle. Then I have showed some other models that performs well but can't beat the score .12324 . Then I have explained why some models with certain parameter works well. After that I showed a table where different models performance is listed and added my analysis and observation. Then I have have Showed four combination of Ensemble and their kaggle score is also attached with them. In the 4th combination of Ensemble method I have found the best kaggle score which is 0.12192. This is the overall best score and achived through combining 4 ann models, xgboost and lasso.

Reference

xgboost:

<https://www.kaggle.com/dansbecker/xgboost> (<https://www.kaggle.com/dansbecker/xgboost>).

<https://medium.com/@gabrieltseng/gradient-boosting-and-xgboost-c306c1bcfaf5> (<https://medium.com/@gabrieltseng/gradient-boosting-and-xgboost-c306c1bcfaf5>).

regression + graph :

<https://www.kaggle.com/janiobachmann/predicting-house-prices-regression-techniques> (<https://www.kaggle.com/janiobachmann/predicting-house-prices-regression-techniques>).

Selecting and Filtering Data

<https://www.kaggle.com/dansbecker/selecting-and-filtering-in-pandas> (<https://www.kaggle.com/dansbecker/selecting-and-filtering-in-pandas>).

Handling Missing Values

<https://www.kaggle.com/dansbecker/handling-missing-values> (<https://www.kaggle.com/dansbecker/handling-missing-values>).

why use conditional probability coding

<https://medium.com/airbnb-engineering/designing-machine-learning-models-7d0048249e69> (<https://medium.com/airbnb-engineering/designing-machine-learning-models-7d0048249e69>).

one hot encoding

<https://hackernoon.com/what-is-one-hot-encoding-why-and-when-do-you-have-to-use-it-e3c6186d008f> (<https://hackernoon.com/what-is-one-hot-encoding-why-and-when-do-you-have-to-use-it-e3c6186d008f>).

<https://medium.com/@rajatgupta310198/getting-started-with-neural-network-for-regression-and-tensorflow-58ad3bd75223> (<https://medium.com/@rajatgupta310198/getting-started-with-neural-network-for-regression-and-tensorflow-58ad3bd75223>).

class example

https://colab.research.google.com/drive/1MExQ52bvHSPaUrGe8RvHZifvE6K6a0qh?fbclid=IwAR2EUWi4q6_q0mFbXQwGh4GNgB2Ex_WpP3K0L12182PdzzsWSsEfzHf0REo#forceEdit=true&offline=true&sandboxMode=true&scrollTo=-Rh3-Vt9Nev9
(https://colab.research.google.com/drive/1MExQ52bvHSPaUrGe8RvHZifvE6K6a0qh?fbclid=IwAR2EUWi4q6_q0mFbXQwGh4GNgB2Ex_WpP3K0L12182PdzzsWSsEfzHf0REo#forceEdit=true&offline=true&sandboxMode=true&scrollTo=-Rh3-Vt9Nev9).

Why cross validation

<https://towardsdatascience.com/5-reasons-why-you-should-use-cross-validation-in-your-data-science-project-8163311a1e79> (<https://towardsdatascience.com/5-reasons-why-you-should-use-cross-validation-in-your-data-science-project-8163311a1e79>)

Decision Tree - Regression

https://www.saedsayad.com/decision_tree_reg.htm (https://www.saedsayad.com/decision_tree_reg.htm)

Some more

<https://www.kaggle.com/klyusba/house-prices-advanced-regression-techniques/lasso-model-for-regression-problem/notebook> (<https://www.kaggle.com/klyusba/house-prices-advanced-regression-techniques/lasso-model-for-regression-problem/notebook>)

<https://www.kaggle.com/juliencs/house-prices-advanced-regression-techniques/a-study-on-regression-applied-to-the-ames-dataset/> (<https://www.kaggle.com/juliencs/house-prices-advanced-regression-techniques/a-study-on-regression-applied-to-the-ames-dataset/>)

<https://www.kaggle.com/apapiu/house-prices-advanced-regression-techniques/regularized-linear-models> (<https://www.kaggle.com/apapiu/house-prices-advanced-regression-techniques/regularized-linear-models>)

<https://www.kaggle.com/juliencs/a-study-on-regression-applied-to-the-ames-dataset> (<https://www.kaggle.com/juliencs/a-study-on-regression-applied-to-the-ames-dataset>)

For descriptive section

I have inspired from Ian Goodfellow's book and used his way of explanation to explain my choice. His book can be found here: <https://www.deeplearningbook.org/> (<https://www.deeplearningbook.org/>)